

DHCP Snooping

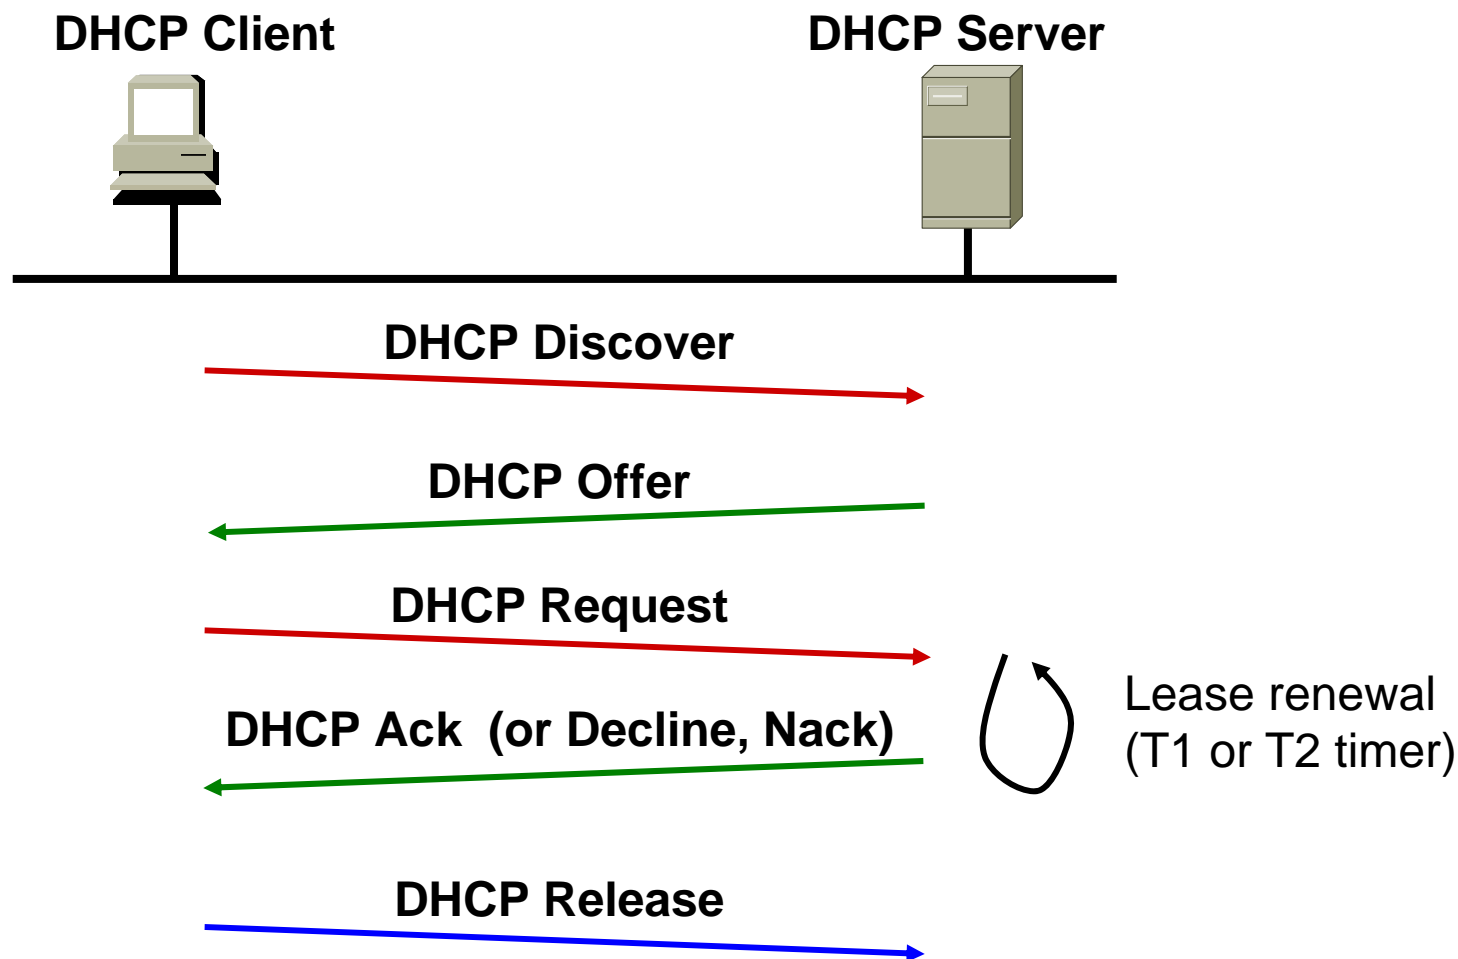
Mini Primer on DHCP (RFC 2131 and 2132)

Cisco.com

- Centralized administration of IP address config
- Superset of BootP Client/Server protocol
- Temporary allocation of IP address and options based on MAC, Client ID, or subnet (GIADDR)
- Transport: UDP, port 67 (server listens on this port) and 68 (client listens on this port)
- Lease renewal efforts occur at two intervals:
 - T1 – 1/2 of the lease has been used
 - T2 – 7/8 of the lease has been used

DHCP Address Acquisition

Cisco.com



DHCP Discover (client-to-server)

No.	Time	Source	Destination	Protocol	Info
37	15.956730	0.0.0.0	255.255.255.255	DHCP	DHCP Discover - Transaction ID 0x7dab8029
38	15.962968	171.69.80.3	171.69.81.43	DHCP	DHCP Offer - Transaction ID 0x7dab8029
39	15.963357	0.0.0.0	255.255.255.255	DHCP	DHCP Request - Transaction ID 0x7dab8029
40	15.968051	171.69.80.2	171.69.81.43	DHCP	DHCP Offer - Transaction ID 0x7dab8029
41	15.973204	171.69.80.2	171.69.81.43	DHCP	DHCP ACK - Transaction ID 0x7dab8029
42	15.981081	171.69.80.3	171.69.81.43	DHCP	DHCP ACK - Transaction ID 0x7dab8029

+

Frame 37 (342 bytes on wire, 342 bytes captured)

+

Ethernet II, Src: Foxconn_ea:6a:d8 (00:01:6c:ea:6a:d8), Dst: Broadcast (ff:ff:ff:ff:ff:ff)

+

Internet Protocol, Src: 0.0.0.0 (0.0.0.0), Dst: 255.255.255.255 (255.255.255.255)

+

User Datagram Protocol, Src Port: bootpc (68), Dst Port: bootps (67)

+

Bootstrap Protocol

Message type: Boot Request (1)

Hardware type: Ethernet

Hardware address length: 6

Hops: 0

Transaction ID: 0x7dab8029

Seconds elapsed: 0

Bootp flags: 0x0000 (Unicast)

Client IP address: 0.0.0.0 (0.0.0.0)

Your (client) IP address: 0.0.0.0 (0.0.0.0)

Next server IP address: 0.0.0.0 (0.0.0.0)

Relay agent IP address: 0.0.0.0 (0.0.0.0)

Client MAC address: Foxconn_ea:6a:d8 (00:01:6c:ea:6a:d8)

Server host name not given

Boot file name not given

Magic cookie: (OK)

Option 53: DHCP Message Type = DHCP Discover

Option 116: DHCP Auto-Configuration (1 bytes)

Option 61: Client identifier

Option 50: Requested IP Address = 1.1.1.1

Option 12: Host Name = "kbogart-wxp04"

Option 60: vendor class identifier = "MSFT 5.0"

Option 55: Parameter Request List

End option

This is the last IP Address I had. In this case it was a static IP address I assigned before I changed it over to "dynamic" to force the DHCP process.

DHCP Offer (server-to-client)

No.	Time	Source	Destination	Protocol	Info
37	15.956730	0.0.0.0	255.255.255.255	DHCP	DHCP Discover - Transaction ID 0x7dab8029
38	15.962968	171.69.80.3	171.69.81.43	DHCP	DHCP offer - Transaction ID 0x7dab8029
39	15.963357	0.0.0.0	255.255.255.255	DHCP	DHCP Request - Transaction ID 0x7dab8029
40	15.968051	171.69.80.2	171.69.81.43	DHCP	DHCP offer - Transaction ID 0x7dab8029
41	15.973204	171.69.80.2	171.69.81.43	DHCP	DHCP ACK - Transaction ID 0x7dab8029
42	15.981081	171.69.80.3	171.69.81.43	DHCP	DHCP ACK - Transaction ID 0x7dab8029

+	Frame 38 (351 bytes on wire, 351 bytes captured)
+	Ethernet II, Src: 00:19:07:ba:7d:c0 (00:19:07:ba:7d:c0), Dst: Foxconn_ea:6a:d8 (00:01:6c:ea:6a:d8)
+	Internet Protocol, Src: 171.69.80.3 (171.69.80.3), Dst: 171.69.81.43 (171.69.81.43)
+	User Datagram Protocol, Src Port: bootps (67), Dst Port: bootpc (68)
-	Bootstrap Protocol
	Message type: Boot Reply (2)
	Hardware type: Ethernet
	Hardware address length: 6
	Hops: 0
	Transaction ID: 0x7dab8029
	Seconds elapsed: 0
+	Bootp flags: 0x0000 (Unicast)
	Client IP address: 0.0.0.0 (0.0.0.0)
	Your (client) IP address: 171.69.81.43 (171.69.81.43)
	Next server IP address: 171.71.179.41 (171.71.179.41)
	Relay agent IP address: 171.69.80.3 (171.69.80.3)
	Client MAC address: Foxconn_ea:6a:d8 (00:01:6c:ea:6a:d8)
	Server host name not given
	Boot file name: /x86pc/undi/bstrap/bstrap.0
	Magic cookie: (OK)
	Option 53: DHCP Message Type = DHCP offer
	Option 54: Server Identifier = 171.68.10.69
	Option 51: IP Address Lease Time = 2 days, 18 hours, 40 minutes, 35 seconds
	Option 1: Subnet Mask = 255.255.254.0
	Option 15: Domain Name = "cisco.com"
	Option 3: Router = 171.69.80.1
+	Option 6: Domain Name Server
+	Option 44: NetBIOS over TCP/IP Name Server
+	Option 46: NetBIOS over TCP/IP Node Type = 1 (Broadcast)

Downloaded from <http://ajphaphysocpharm.sagepub.com/> at 11:00 11 September 2014

No.	Time	Source	Destination	Protocol	Info
37	15.956730	0.0.0.0	255.255.255.255	DHCP	DHCP Discover - Transaction ID 0x7dab8029
38	15.962968	171.69.80.3	171.69.81.43	DHCP	DHCP Offer - Transaction ID 0x7dab8029
39	15.963357	0.0.0.0	255.255.255.255	DHCP	DHCP Request - Transaction ID 0x7dab8029
40	15.968051	171.69.80.2	171.69.81.43	DHCP	DHCP Offer - Transaction ID 0x7dab8029
41	15.973204	171.69.80.2	171.69.81.43	DHCP	DHCP ACK - Transaction ID 0x7dab8029
42	15.981081	171.69.80.3	171.69.81.43	DHCP	DHCP ACK - Transaction ID 0x7dab8029

Frame 39 (378 bytes on wire, 378 bytes captured)

⊕ Ethernet II, Src: Foxconn_ea:6a:d8 (00:01:6c:ea:6a:d8), Dst: Broadcast (ff:ff:ff:ff:ff:ff)

Internet Protocol, Src: 0.0.0.0 (0.0.0.0), Dst: 255.255.255.255 (255.255.255.255)

⊕ User Datagram Protocol, Src Port: bootpc (68), Dst Port: bootps (67)

[- Bootstrap Protocol]

Message type: Boot Request (1)

```
Hardware type: Ethernet
```

```
Hardware address length: 6
```

Hops: 0

Transaction ID: 0x7dab8029

```
Seconds elapsed: 0
```

+ Bootp flags: 0x0000 (Unicast)

```
Client IP address: 0.0.0.0 (0.0.0.0)
```

```
Your (client) IP address: 0.0.0.0 (0.0.0.0)
```

```
Next server IP address: 0.0.0.0 (0.0.0.0)
```

Relay agent IP address: 0.0.0.0 (0.0.0.0)

```
Client MAC address: Foxconn_ea:6a:d8 (00:01:6c:ea:6a:d8)
```

```
Server host name not given
```

```

Boot file name not given

```

Magic cookie: (OK)

```
Option 53: DHCP Message Type = DHCP Request
```

```

+ option 61: Client identifier

```

```
Option 50: Requested IP Address = 171.69.81.43
```

```
Option 54: Server Identifier = 171.68.10.69
```

```
Option 12: Host Name = "kbogart-wxp04"
```

```

+ Option 81: FODN

```

```
Option 60: vendor class identifier = "MSFT 5.0"
```

Option 55: Parameter Request List

End Option

Duplicate packets??

- Why do you think my laptop was sent **TWO** DHCP Offers?

37	15.956730	0.0.0.0	255.255.255.255	DHCP	DHCP Discover	- Transaction ID 0x7dab8029
38	15.962968	171.69.80.3	171.69.81.43	DHCP	DHCP Offer	- Transaction ID 0x7dab8029
39	15.963357	0.0.0.0	255.255.255.255	DHCP	DHCP Request	- Transaction ID 0x7dab8029
40	15.968051	171.69.80.2	171.69.81.43	DHCP	DHCP Offer	- Transaction ID 0x7dab8029
41	15.973204	171.69.80.2	171.69.81.43	DHCP	DHCP ACK	- Transaction ID 0x7dab8029
42	15.981081	171.69.80.3	171.69.81.43	DHCP	DHCP ACK	- Transaction ID 0x7dab8029

+	Frame 40 (351 bytes on wire, 351 bytes captured)
+	Ethernet II, Src: 00:19:07:ea:7a:80 (00:19:07:ea:7a:80), Dst: Foxconn_ea:6a:d8 (00:01:6c:ea:6a:d8)
+	Internet Protocol, Src: 171.69.80.2 (171.69.80.2), Dst: 171.69.81.43 (171.69.81.43)
+	User Datagram Protocol, Src Port: bootps (67), Dst Port: bootpc (68)
-	Bootstrap Protocol
	Message type: Boot Reply (2)
	Hardware type: Ethernet
	Hardware address length: 6
	Hops: 0
	Transaction ID: 0x7dab8029
	Seconds elapsed: 0
+	Bootp flags: 0x0000 (Unicast)
	Client IP address: 0.0.0.0 (0.0.0.0)
	Your (client) IP address: 171.69.81.43 (171.69.81.43)
	Next server IP address: 171.71.179.41 (171.71.179.41)
	Relay agent IP address: 171.69.80.2 (171.69.80.2)
	Client MAC address: Foxconn_ea:6a:d8 (00:01:6c:ea:6a:d8)
	Server host name not given
	Boot file name: /x86pc/undi/bstrap/bstrap.0
	Magic cookie: (OK)
	Option 53: DHCP Message Type = DHCP offer
	Option 54: Server Identifier = 171.68.10.69
	Option 51: IP Address Lease Time = 2 days, 18 hours, 40 minutes, 35 seconds
	Option 1: Subnet Mask = 255.255.254.0
	Option 15: Domain Name = "cisco.com"
	Option 3: Router = 171.69.80.1


DHCP ACK (server-to-client)

37	15.956730	0.0.0.0	255.255.255.255	DHCP	DHCP Discover - Transaction ID 0x7dab8029
38	15.962968	171.69.80.3	171.69.81.43	DHCP	DHCP offer - Transaction ID 0x7dab8029
39	15.963357	0.0.0.0	255.255.255.255	DHCP	DHCP Request - Transaction ID 0x7dab8029
40	15.968051	171.69.80.2	171.69.81.43	DHCP	DHCP offer - Transaction ID 0x7dab8029
41	15.973204	171.69.80.2	171.69.81.43	DHCP	DHCP ACK - Transaction ID 0x7dab8029
42	15.981081	171.69.80.3	171.69.81.43	DHCP	DHCP ACK - Transaction ID 0x7dab8029

+	Frame 41 (351 bytes on wire, 351 bytes captured)
+	Ethernet II, Src: 00:19:07:ea:7a:80 (00:19:07:ea:7a:80), Dst: Foxconn_ea:6a:d8 (00:01:6c:ea:6a:d8)
+	Internet Protocol, Src: 171.69.80.2 (171.69.80.2), Dst: 171.69.81.43 (171.69.81.43)
+	User Datagram Protocol, Src Port: bootps (67), Dst Port: bootpc (68)
-	Bootstrap Protocol
	Message type: Boot Reply (2)
	Hardware type: Ethernet
	Hardware address length: 6
	Hops: 0
	Transaction ID: 0x7dab8029
	Seconds elapsed: 0
+	Bootp flags: 0x0000 (Unicast)
	Client IP address: 0.0.0.0 (0.0.0.0)
	Your (client) IP address: 171.69.81.43 (171.69.81.43)
	Next server IP address: 171.71.179.41 (171.71.179.41)
	Relay agent IP address: 171.69.80.2 (171.69.80.2)
	Client MAC address: Foxconn_ea:6a:d8 (00:01:6c:ea:6a:d8)
	Server host name not given
	Boot file name: /x86pc/undi/bstrap/bstrap.0
	Magic cookie: (OK)
	Option 53: DHCP Message Type = DHCP ACK
	Option 54: Server Identifier = 171.68.10.69
	Option 51: IP Address Lease Time = 2 days, 18 hours, 40 minutes, 35 seconds
	Option 1: Subnet Mask = 255.255.254.0
	Option 15: Domain Name = "cisco.com"
	Option 3: Router = 171.69.80.1
+	Option 6: Domain Name Server
+	Option 44: NetBIOS over TCP/IP Name Server
	Option 46: NetBIOS over TCP/IP Node Type = Unicast

Several DHCP message types....

Cisco.com



Client messages:


- Discover
- Request (4 kinds):
 - selecting
 - renew
 - rebind
 - Init/Reboot
- Decline
- Release
- Inform

Server messages:

- Offer
- ACK
- NAK

DHCP Message Format

Cisco.com




0				16				32							
OP Code (1)				HTYPE (1)				HLEN (1)				HOPS (1)			
TRANSACTION ID (4)															
SECONDS (2)								UNUSED (2)							
CLIENT IP ADDRESS (4)															
YOUR IP ADDRESS (4)															
SERVER IP ADDRESS (4)															
GATEWAY IP ADDRESS (GiADDR) (4)															
SERVER HOST NAME (64)															
BOOT FILE NAME (128)															
VENDOR-SPECIFIC OPTIONS (312)															

The GIADDR is "stuffed" with IP address by IP Helper feature to ID subnet of client

DHCP Spoofing Attack

Cisco.com



Who:

- Malicious user: pretend to be the network DHCP server
- Mis-configured user: fire up DHCP server incorrectly

Where:

- Commonly seen in higher education, metro Ethernet

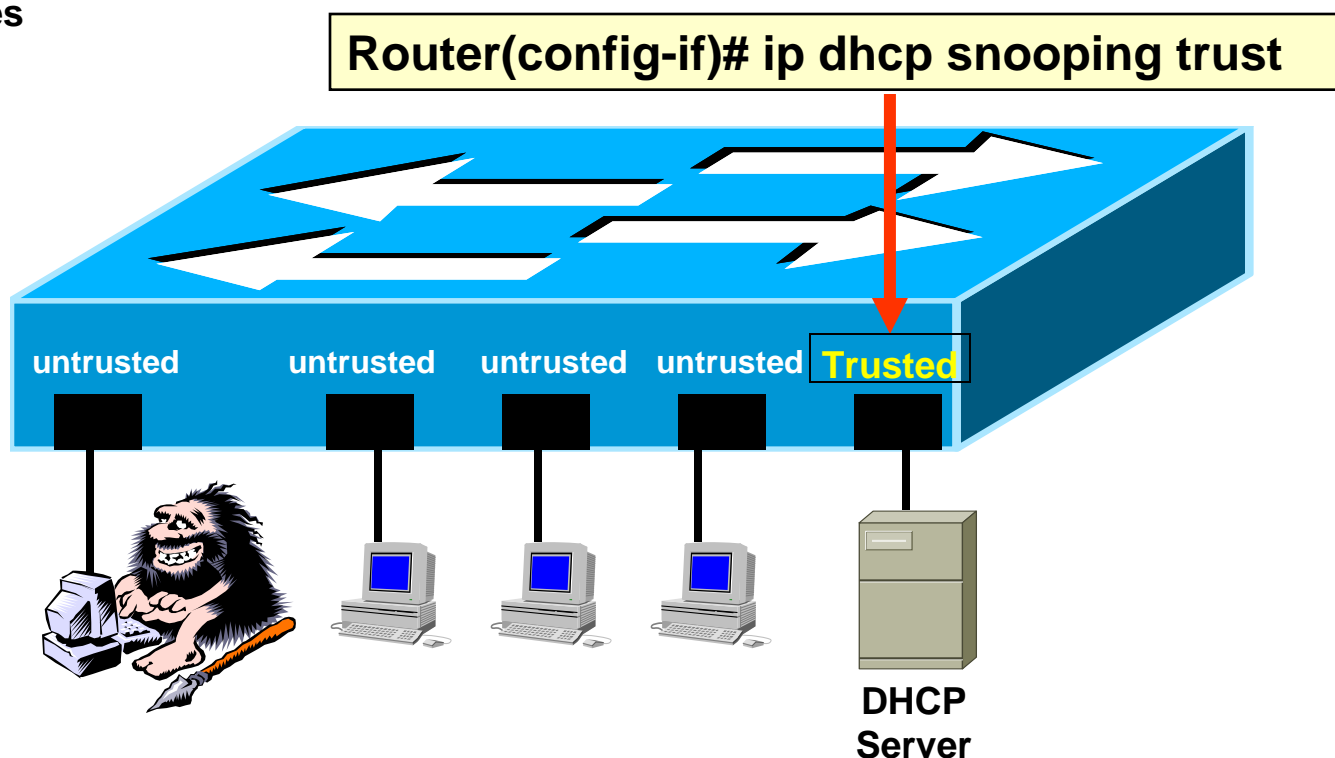
How:

- Attacker Intercepts Discovery Broadcast and Replies With Bogus Gateway and DNS Addresses

Do I Trust You?

Cisco.com

- DHCP Snooping relies on correct identification of Trusted and Untrusted ports.
- Default = **All Ports Untrusted**
- Trust ONLY those ports for which you have direct control of the end-device, ie:
 - ✓Routers
 - ✓Switches
 - ✓Servers




DHCP Attack Solution: DHCP Snooping

Cisco.com

DHCP Snooping – discarding attacker’s bogus DHCP offer messages by intercepting DHCP messages within a switch

- Switch forwards DHCP requests from untrusted access ports only to Trusted ports.
- All other types of DHCP traffic from untrusted access ports dropped.
- If network DHCP server not local to the switch, trust the uplink port
- **Building a DHCP binding table containing client IP address, client MAC address, port, VLAN number...**
- Optional insertion and removal of DHCP option 82 data into/from DHCP messages
- DoS attack on DHCP server is prevented by rate limiting DHCP packets per access port

DHCP Snooping



DHCP Binding Table


Cisco.com

- Contains binding entries for local **untrusted** ports only
- Includes both static entries and dynamic entries learned via DHCP gleaning

4 bytes	IP Address
6 bytes	MAC Address
2 bytes	VLAN Id
4 bytes	Lease Timer
4 bytes	Port
4 bytes	Binding Type


What is allowed to pass (client-to-server)?

Cisco.com



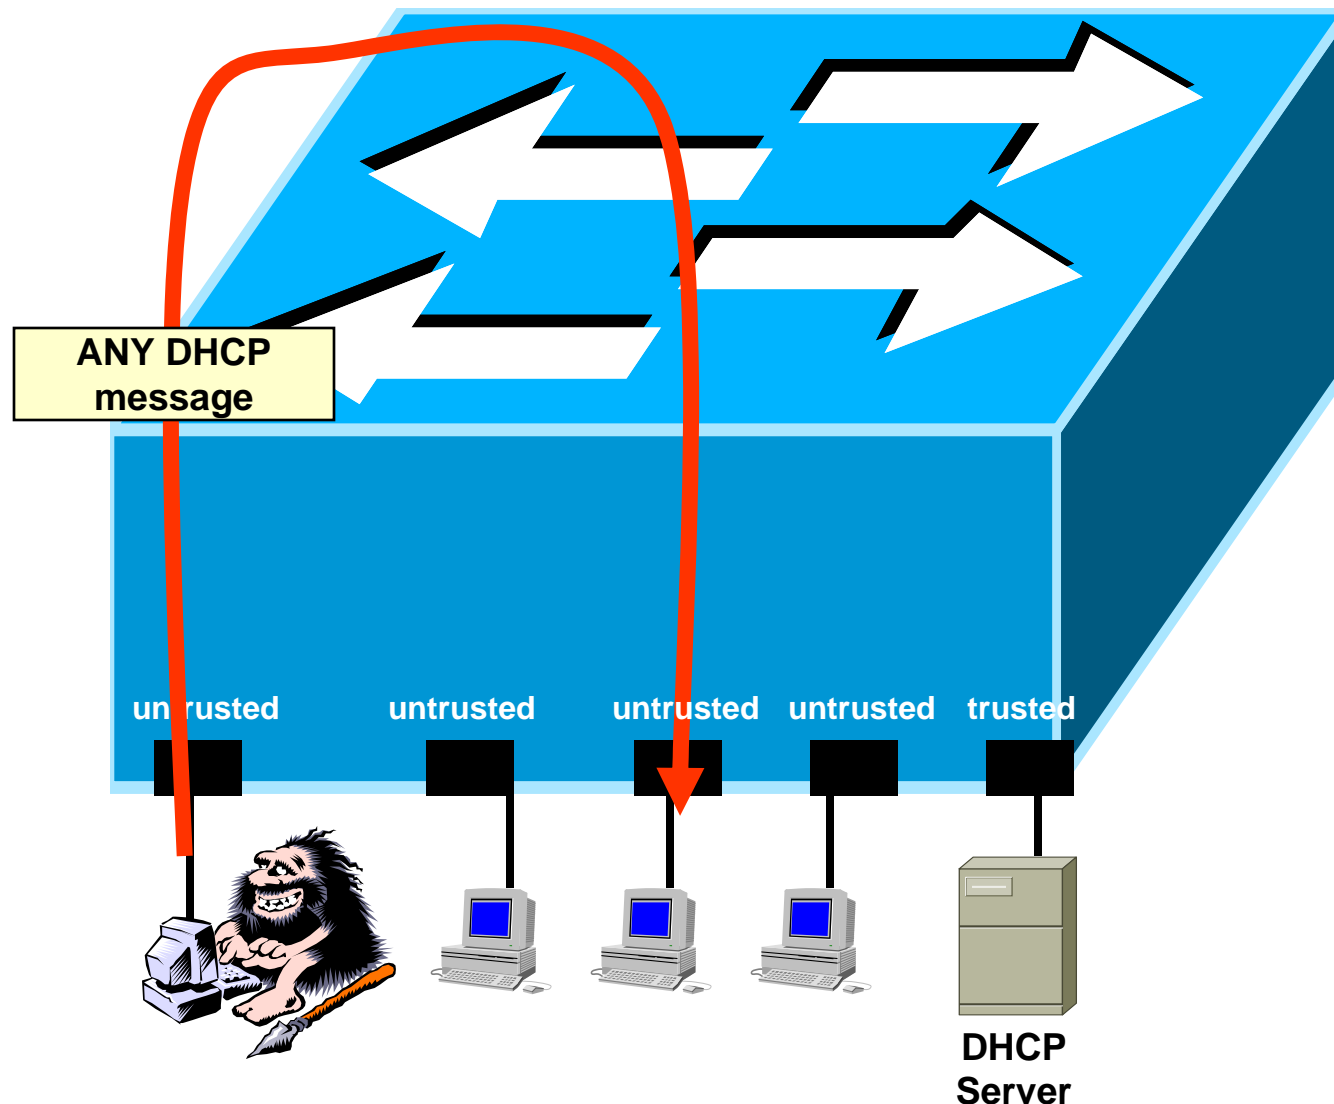
What is allowed to pass (server-to-client)?

Cisco.com



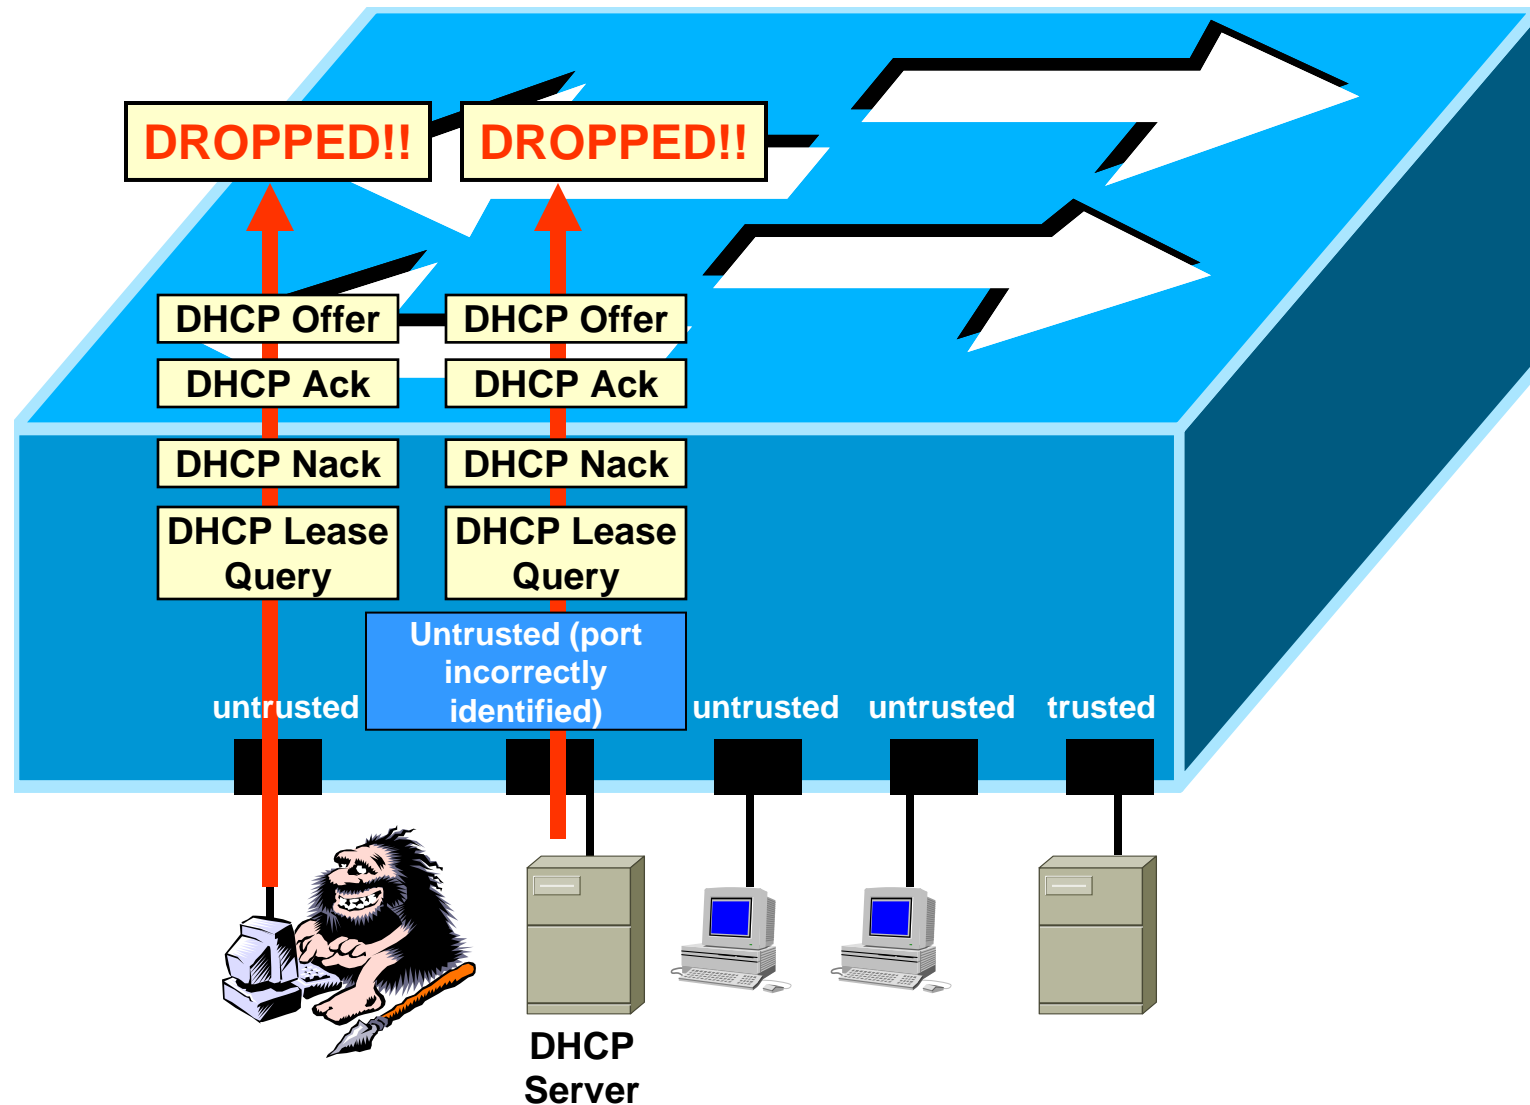
What is prevented (**untrusted-to-untrusted**).

Cisco.com



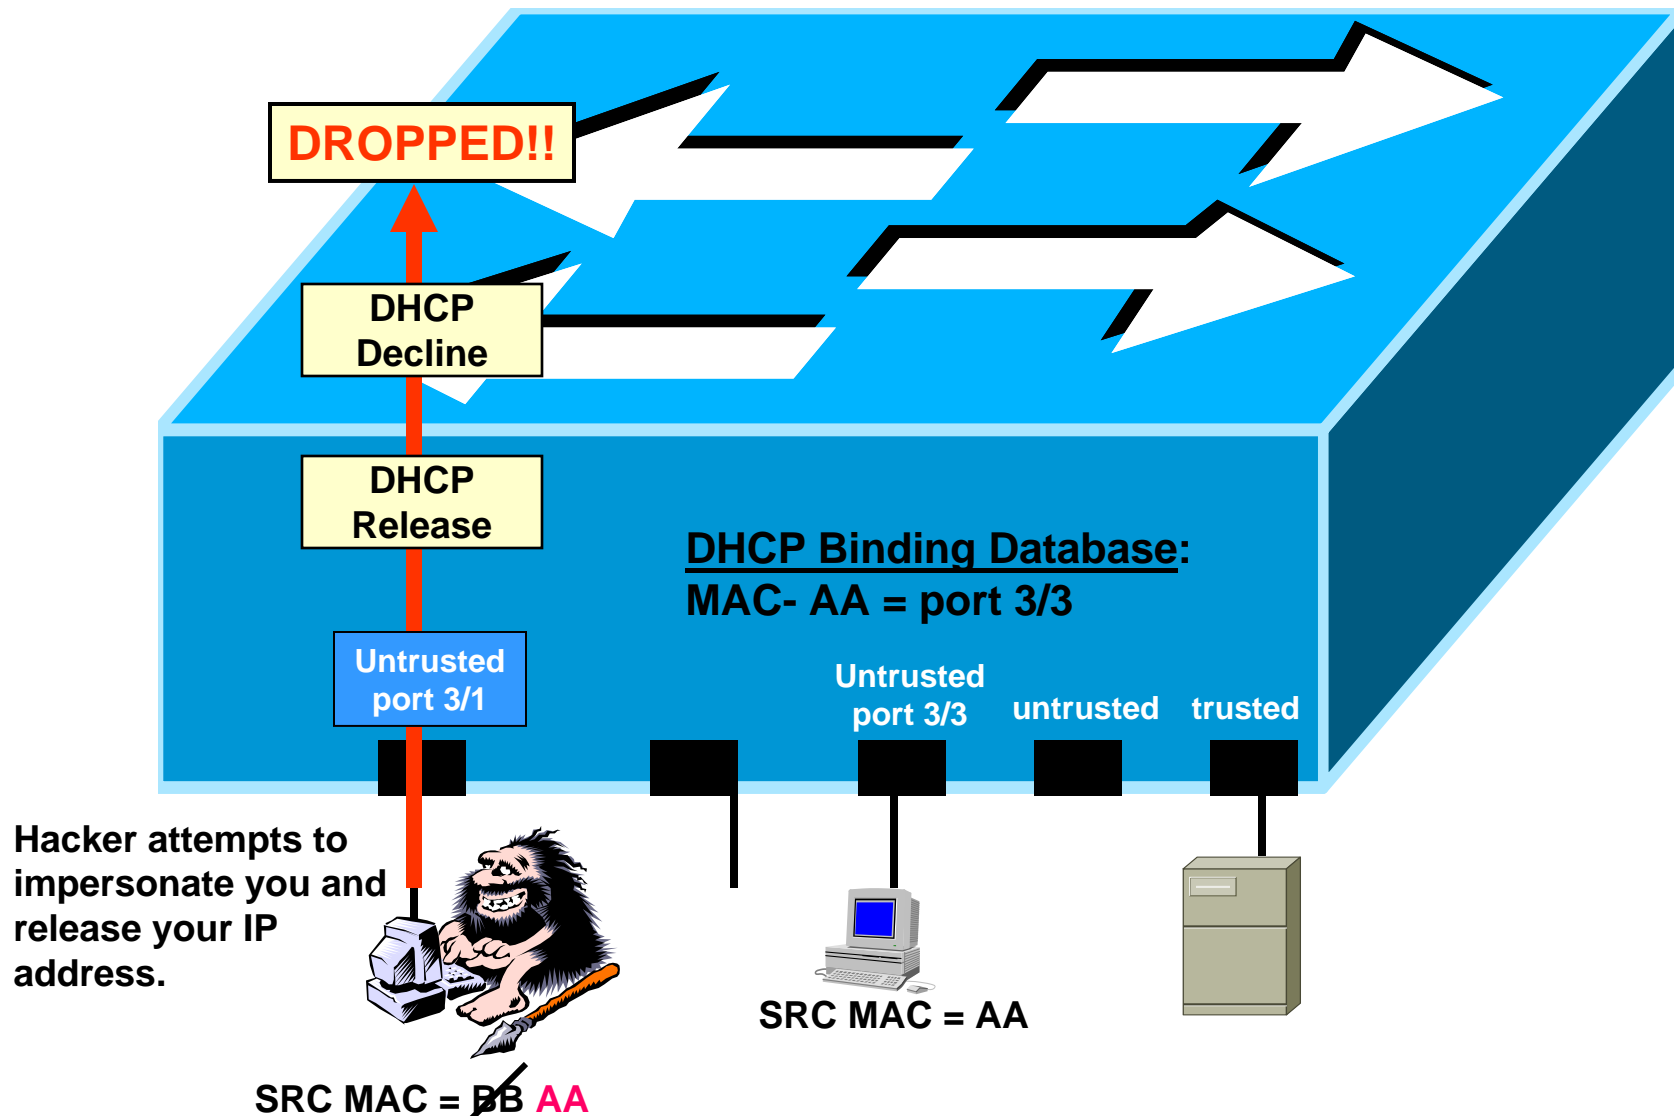
What is prevented (**Untrusted Server Packets**).

Cisco.com



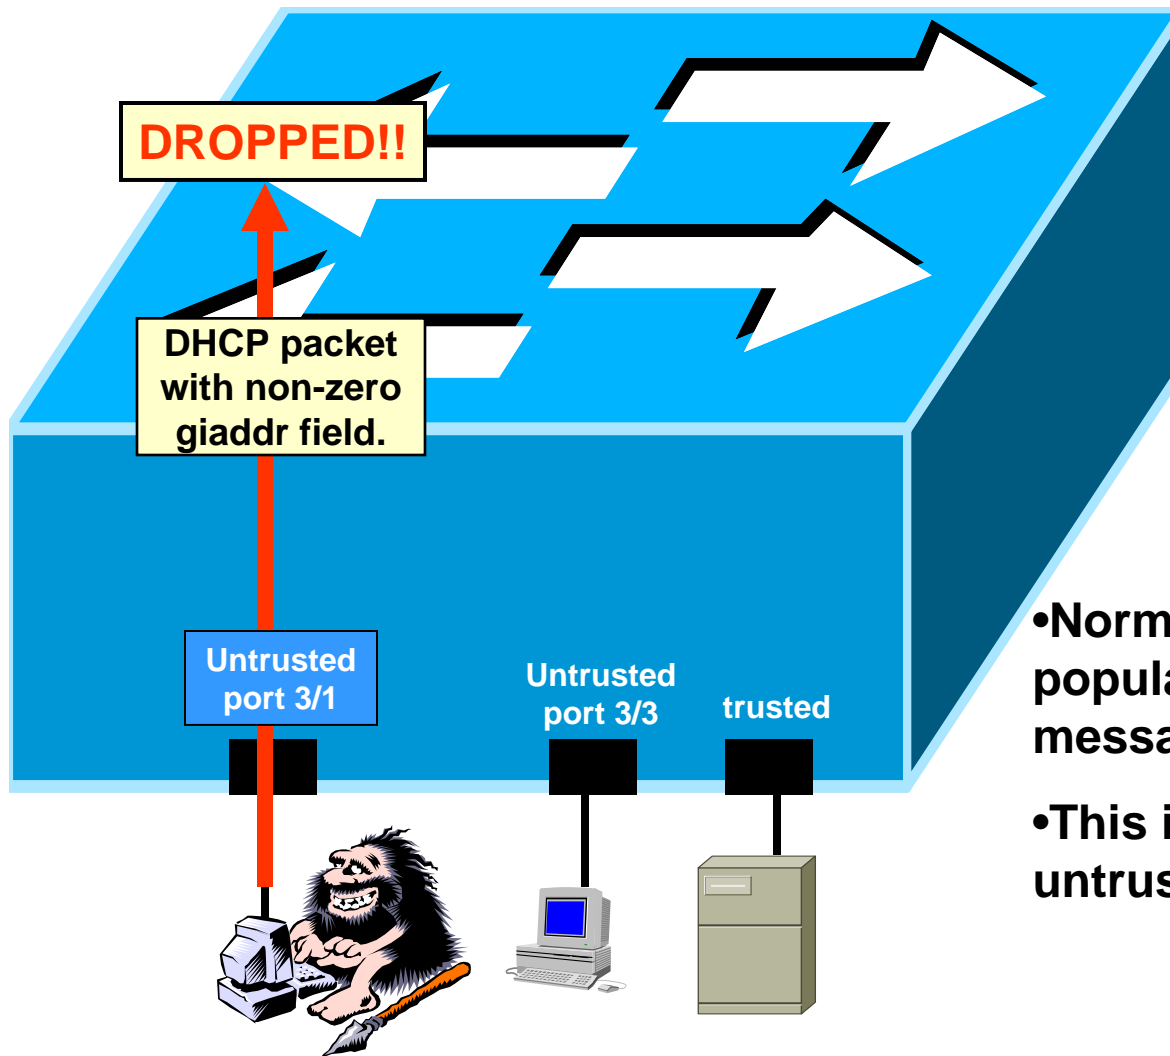
What is prevented (**Who do you think YOU are??**).

Cisco.com



What is prevented (**No relay for YOU!!**).

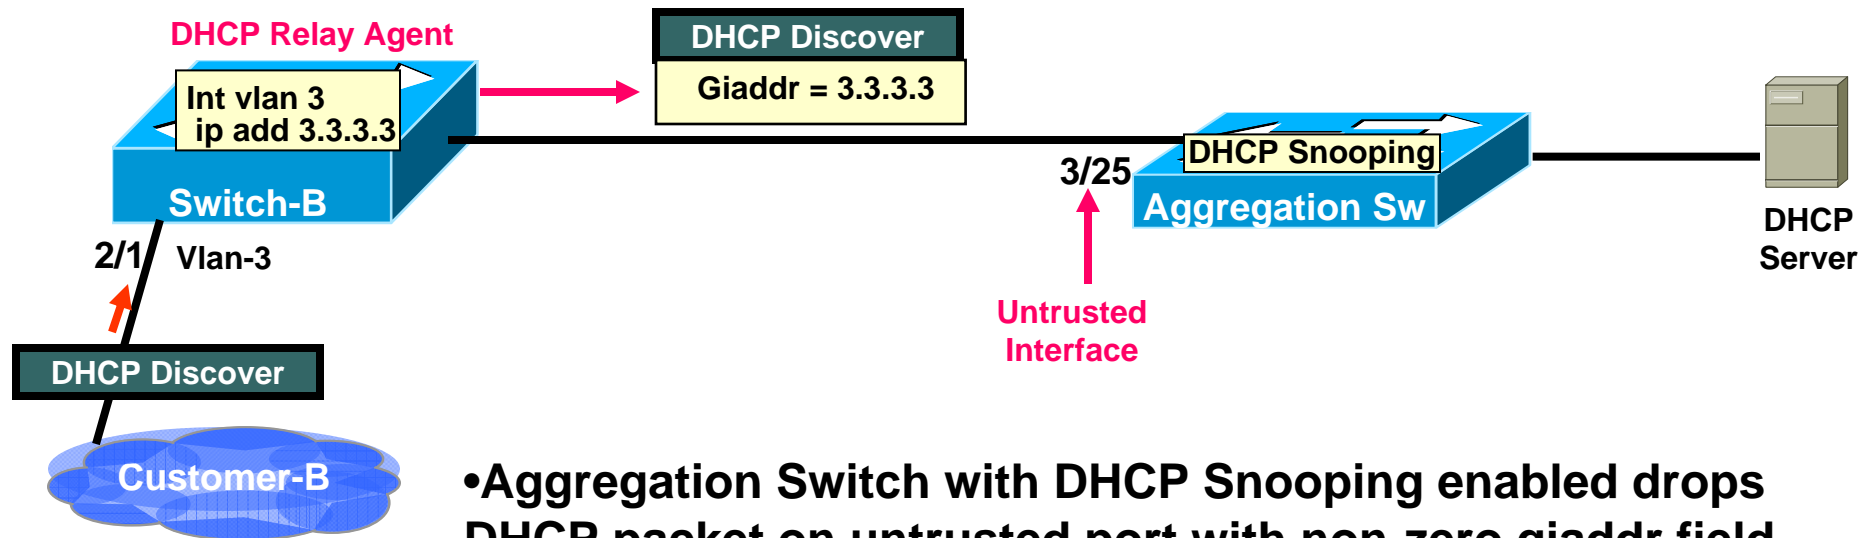
Cisco.com



- Normal DHCP-Relay operation populates “giaddr” field in DHCP messages.
- This is not allowed if arriving on an untrusted port.

DHCP Relay packet dropped!!

Cisco.com



- Aggregation Switch with DHCP Snooping enabled drops DHCP packet on untrusted port with non-zero giaddr field.

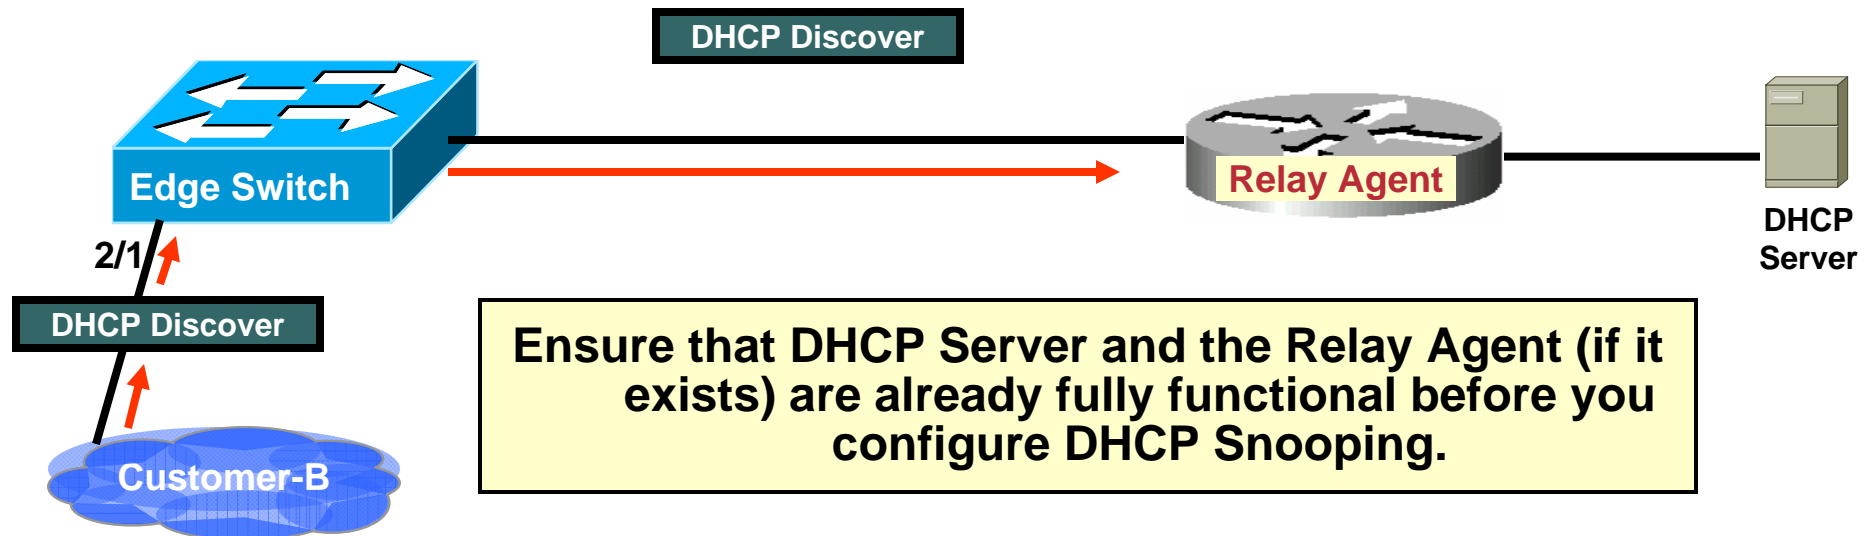
```
Mar 20 12:06:03: %DHCP_SNOOPING-5-DHCP_SNOOPING_NONZERO_GIADDR: DHCP_SNOOPING drop
message with non-zero giaddr or option82 value on untrusted port, message type: DHC
PDISCOVER, MAC sa: 0013.5f1d.7f80
```

The Solution:

```
Cat4500(config)#int fast 3/25
Cat4500(config-if)#ip dhcp snooping trust
Cat4500(config-if)#end
Cat4500#
```

DHCP Snooping - Configuration

Cisco.com



```
Cat3750(config)#
Cat3750(config)#
Cat3750(config)#ip dhcp snooping
Cat3750(config)#ip dhcp snooping vlan 2-8
Cat3750(config)#interface gig 1/0/2
Cat3750(config-if)#no ip dhcp snooping trust
Cat3750(config-if)#ip dhcp snooping trust
Cat3750(config-if)#end
Cat3750#
```

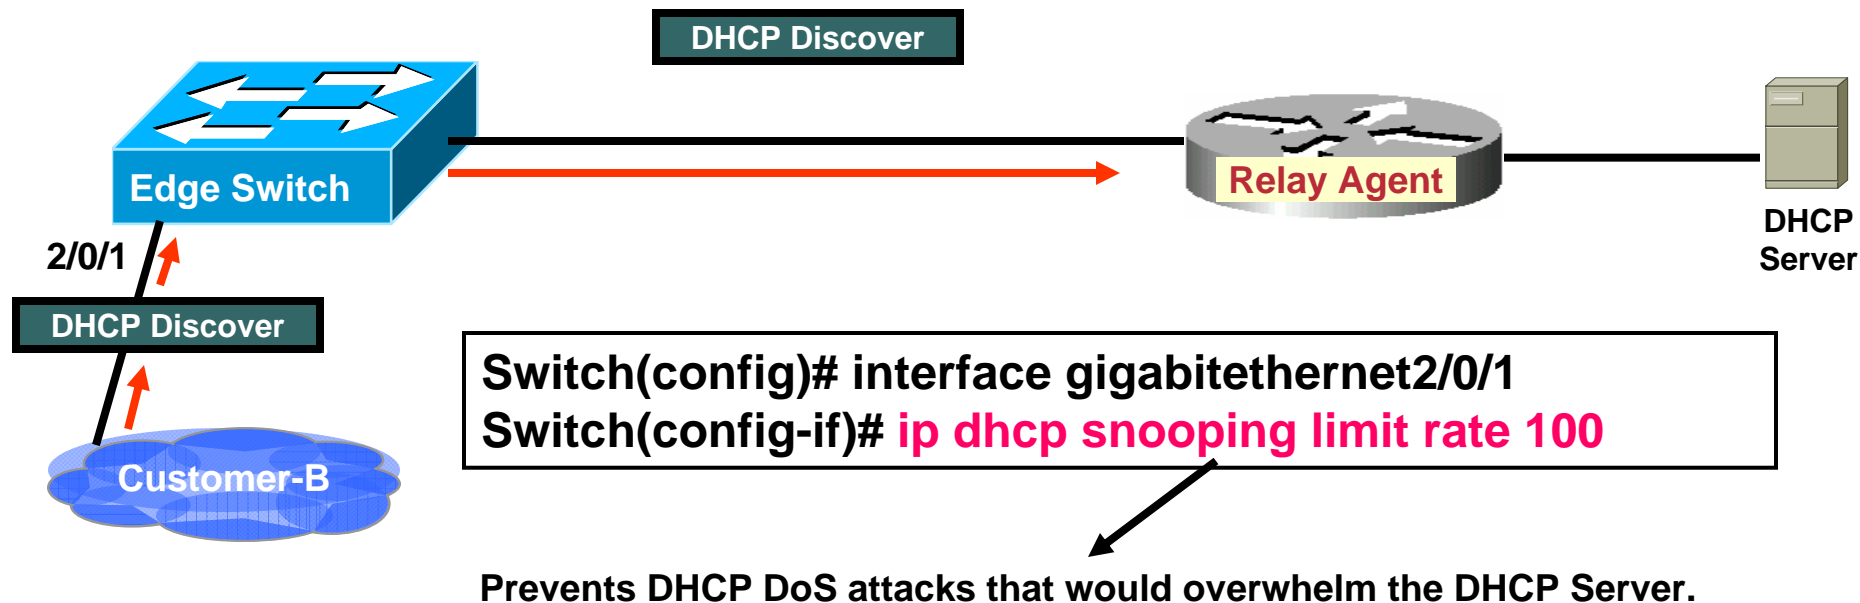
Enables DHCP Snooping only for these VLANs

Default setting

Configure this on ports leading to trusted DHCP Servers...or on uplink ports to Aggregation Switches.

DHCP Snooping – Additional Config Options

Cisco.com



- DHCP Snooping can also be configured on Private VLANs.
- Must configure **only on the Primary VLAN**...will be dynamically propagated to all Secondary VLANs.
- No way (currently) to have different DHCP Snooping configurations applied to Secondary VLANs all residing under the same Primary VLAN.

DHCP Snooping – Verification

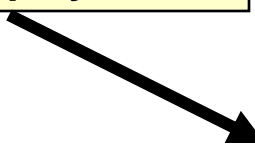
Cisco.com

```
Cat6500#sho ip dhcp snooping
Switch DHCP snooping is enabled
DHCP snooping is configured on following VLANs:
1-12
DHCP snooping is operational on following VLANs:
1-8,10,12
DHCP snooping is configured on the following L3 Interfaces:

Insertion of option 82 is enabled
Option 82 on untrusted port is not allowed
Verification of hwaddr field is enabled
Verification of giaddr field is enabled
DHCP snooping trust/rate is configured on the following Interfaces:

Interface                               Trusted      Rate limit (pps)
-----
FastEthernet3/7                        yes          unlimited
Cat6500#
```

**Untrusted
interfaces don't
display.**



```
Cat6500#sho ip dhcp snooping bind
MacAddress      IpAddress      Lease(sec)      Type            VLAN  Interface
-----
00:11:5C:16:4F:60  1.1.1.1        215944          dhcp-snooping   1     FastEthernet3/6
Total number of bindings: 1

Cat6500#
```


DHCP Relay Agent

Cisco.com

- **Best practice is to store DHCP Binding Database externally to the switch.**
 - If stored locally in flash/bootflash, database must be erased and re-written for every new entry.
 - CPU intensive...can lock up the switch.
 - If switch crashes or reloads, all entries / lease info lost and can kill the DHCP Snooping process.
- **Feature to do this is called “DHCP Snooping Database Agent”.**

Can also use FTP, HTTP, and RCP

Switch(Config)# ip dhcp snooping database tftp://192.168.1.1/Snoop-data.dhcp
Switch(Config)# ip dhcp snooping database write-delay 15

Specify the duration for which the transfer should be delayed after the binding database changes. The range is from 15 to 86400 seconds. The default is 300 seconds (5 minutes).

DHCP Relay Agent GOTCHAS (1)

Cisco.com

From the Cat3750 Configuration Guide:

- “For network-based URLs (such as TFTP and FTP), you must create an empty file at the configured URL before the switch can write bindings to the binding file at that URL. See the documentation for your TFTP server to determine whether you must first create an empty file on the server; some TFTP servers cannot be configured this way.”

Meaning – The switch cannot create this file from scratch. The server must already contain a 0-byte file with this name for this to work.

- What will you see if you DON'T have a 0-byte file to start with??

```
Cat3750# show ip dhcp snooping database
```

```
Agent URL : tftp://192.168.1.1/Snoop-data.dhcp
```

```
Agent Running : No
```

```
Delay Timer Expiry : Not Running
```

```
Abort Timer Expiry : Not Running
```

```
Last Succeeded Time : None
```

```
Last Failed Time : 18:56:49 DST Mon Sep 18 2006
```

```
Last Failed Reason : New line expected in database.
```

Total Attempts	:	9285	Startup Failures	:	9284
Successful Transfers	:	0	Failed Transfers	:	9285
Successful Reads	:	0	Failed Reads	:	1
Successful Writes	:	0	Failed Writes	:	0
Media Failures	:	0			

DHCP Relay Agent GOTCHAS (2)

Cisco.com

From the Cat3750 Configuration Guide:

- “To ensure that the lease time in the database is accurate, **we recommend** that you enable and configure NTP.”

The REAL STORY: **NTP (Network Time Protocol) is MANDATORY!** Agent won't work without it!!

- What will you see if you DON'T have NTP running??

```
*Jul 27 23:08:20: Safe write timer expired.  
*Jul 27 23:08:20: Trying to open url in safe write mode..  
*Jul 27 23:08:20: Safe write mode failed. Restarting timer.
```

From Case# 601706547: “Safe read write mode is a special mode which tries to open the file mentioned in the database URL in the read-only format and only if it exists tries to write to it as in try to update it. If the file does not exist , it tries to create the file. And from the debug messages (the explanation of debug messages are not documented on CCO) , what I can see is safe mode is just simply failing to access the file.

- A sniffer trace from the customer showed that the switch wasn't sending **ANYTHING** to the database server (which contradicted the explanation above because it wasn't even trying to read the file).
- Turned out that the NTP server had failed, which caused this problem.

I need an NTP Server??!!

Cisco.com

- If the customer doesn't normally use NTP (but it's required for the Database Agent) simply configure the NTP Server on another networking device.

NTP Client (DHCP Snooping Switch)

Cat6500#sh run

Building configuration...

service timestamps debug datetime msec localtime

service timestamps log datetime msec localtime

clock timezone PST -8

clock summer-time PDT recurring

clock calendar-valid

!

interface Vlan1

ip address 1.1.1.12 255.255.255.0

ip helper-address 12.12.12.6

!

ntp logging

ntp clock-period 17179871 --> This is inserted by default, no need to modify

ntp source Vlan1

ntp update-calendar

ntp server 1.1.1.10

!

end

NTP Server (Any IOS Router/Switch with IP connectivity to the NTP Client)

Cat4500#sh run

Building configuration...

service timestamps debug uptime

service timestamps log datetime

!

hostname Cat4500

!

interface Vlan1

ip address 1.1.1.10 255.255.255.0

!

ntp source Vlan1

ntp master 1

ntp update-calendar

ntp peer 1.1.1.12 source Vlan1

end

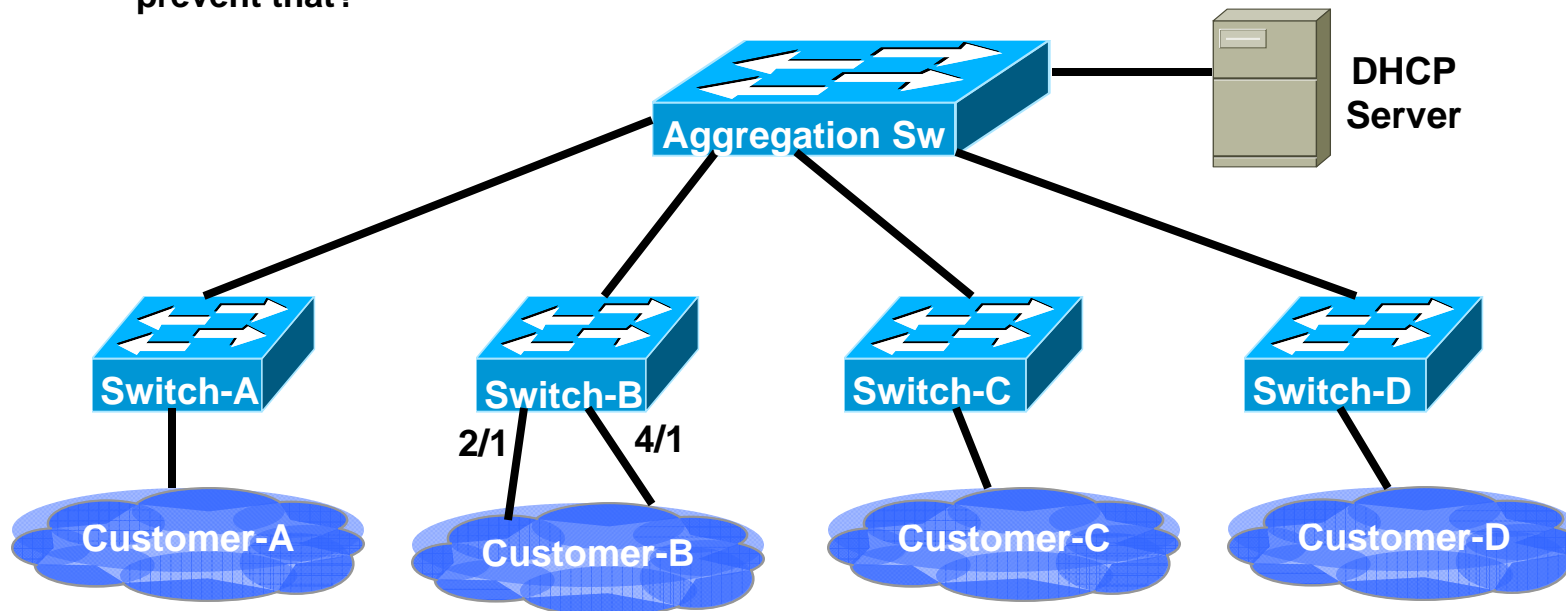
Cat4500#**clock set 19:01:30 19 March 2008 --> Don't forget to set the clock!!**

Restricting Allocated Addresses

Cisco.com

Customer's Challenge:

1. "How can I ensure that each switch is only allocated a maximum of "X" addresses from my DHCP Pool?"
2. "How can I ensure that port 2/1 on Switch-B is only allocated a maximum of "X" addresses from my DHCP Pool?"
3. "What if someone in Customer-C's network is attempting a DHCP DoS attack (sending multiple DHCPDiscover/Request messages to completely exhaust the DHCP Address Pool)? How can I prevent that?"

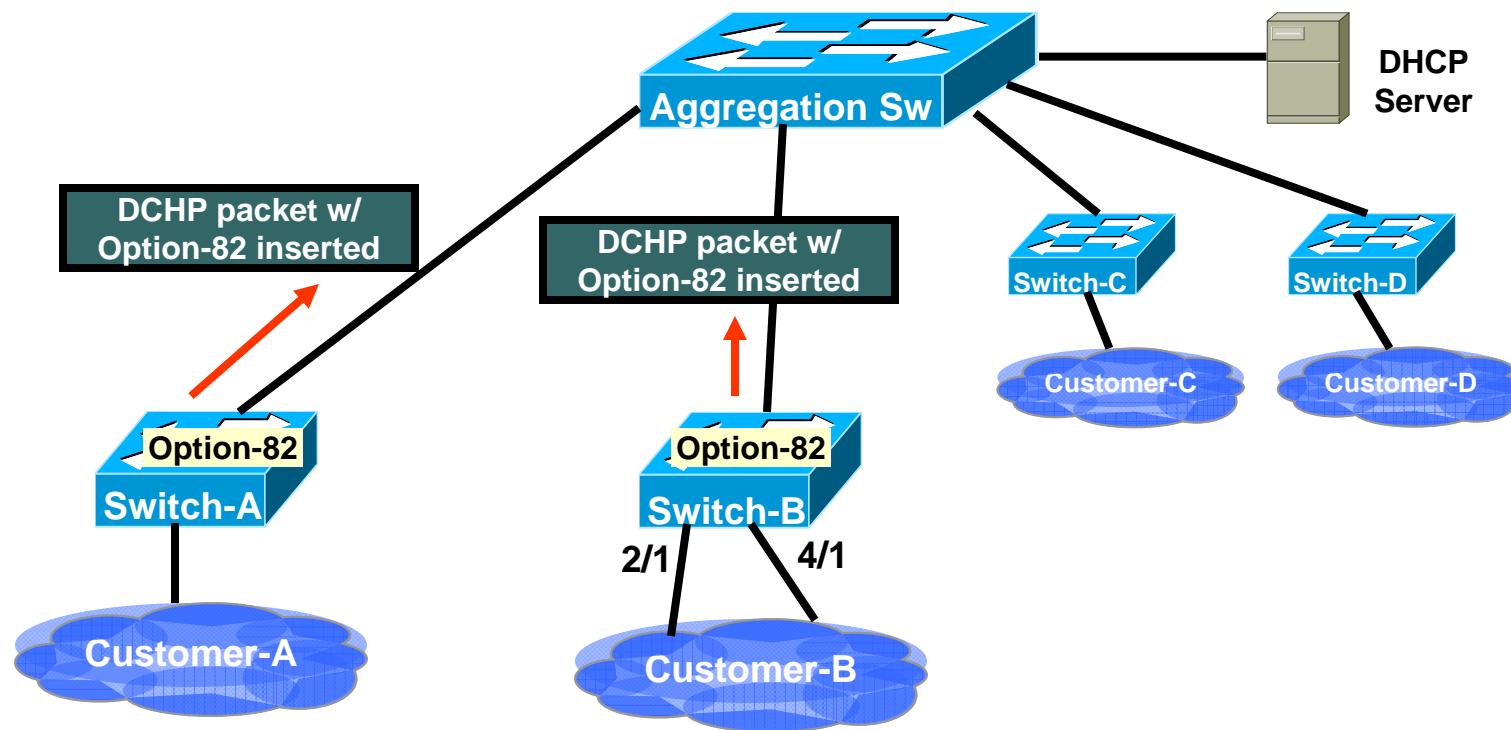


The Solution: DHCP Option-82

a.k.a. DHCP Relay Agent Option (RFC 3046)

DHCP Option-82

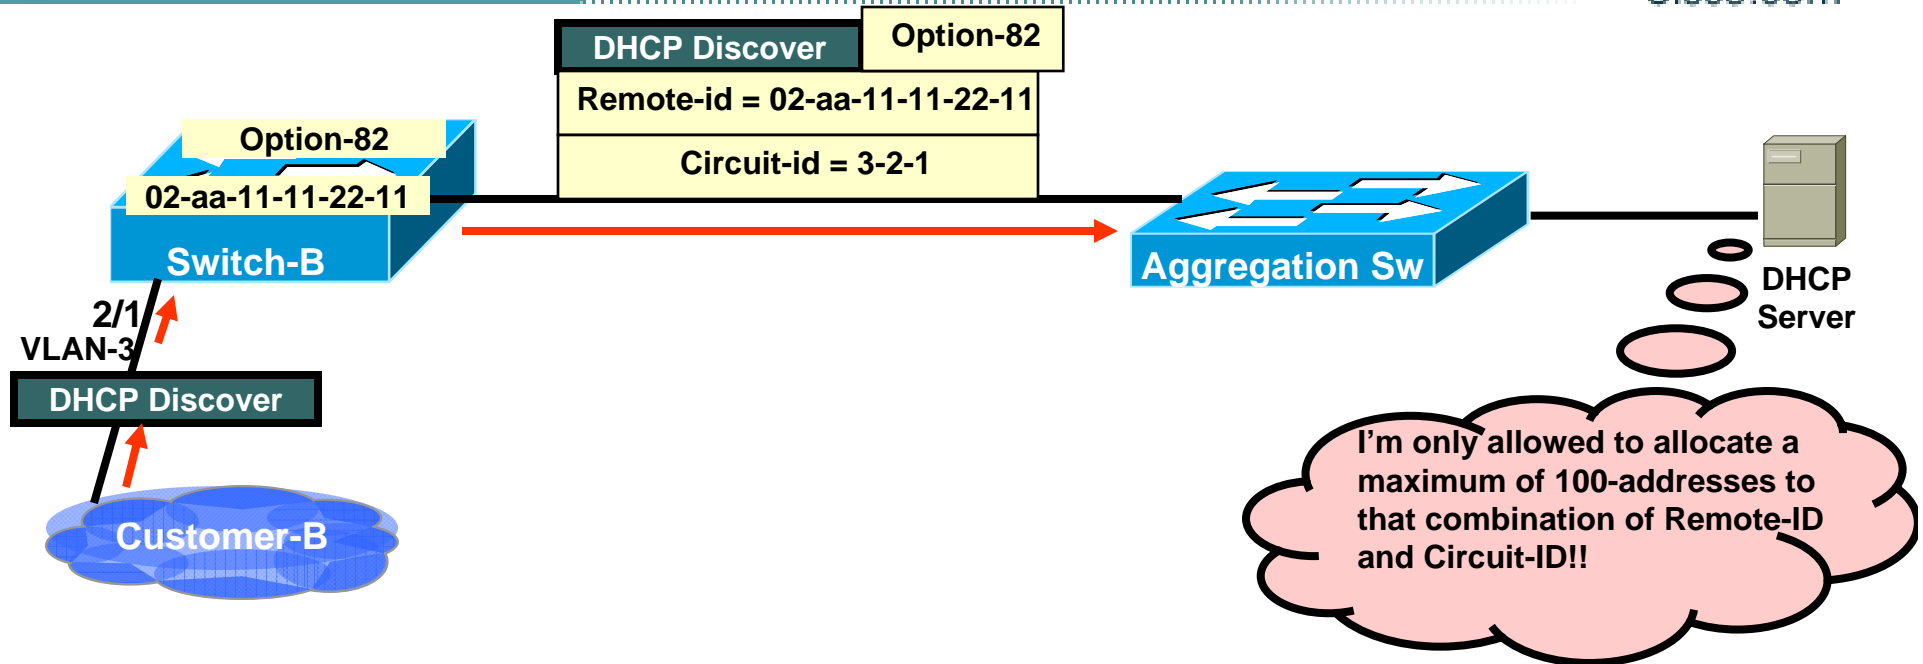
Cisco.com



1. Option-82 allows trusted access devices to insert this option into (and remove from) DHCP Packets.
2. This option gives descriptive information about the device/port that received the DHCP message.

DHCP Option-82

Cisco.com




1. Switch adds "Remote-ID" and "Circuit-ID" sub-options into Option-82 data.
 - Remote-ID default is switch MAC address
 - Circuit-ID default is port identifier in the format "vlan-mod-port"
2. These fields are configurable to use ASCII strings if you prefer


DHCP Option-82 – Technical Details

Cisco.com

Circuit ID Suboption Frame Format



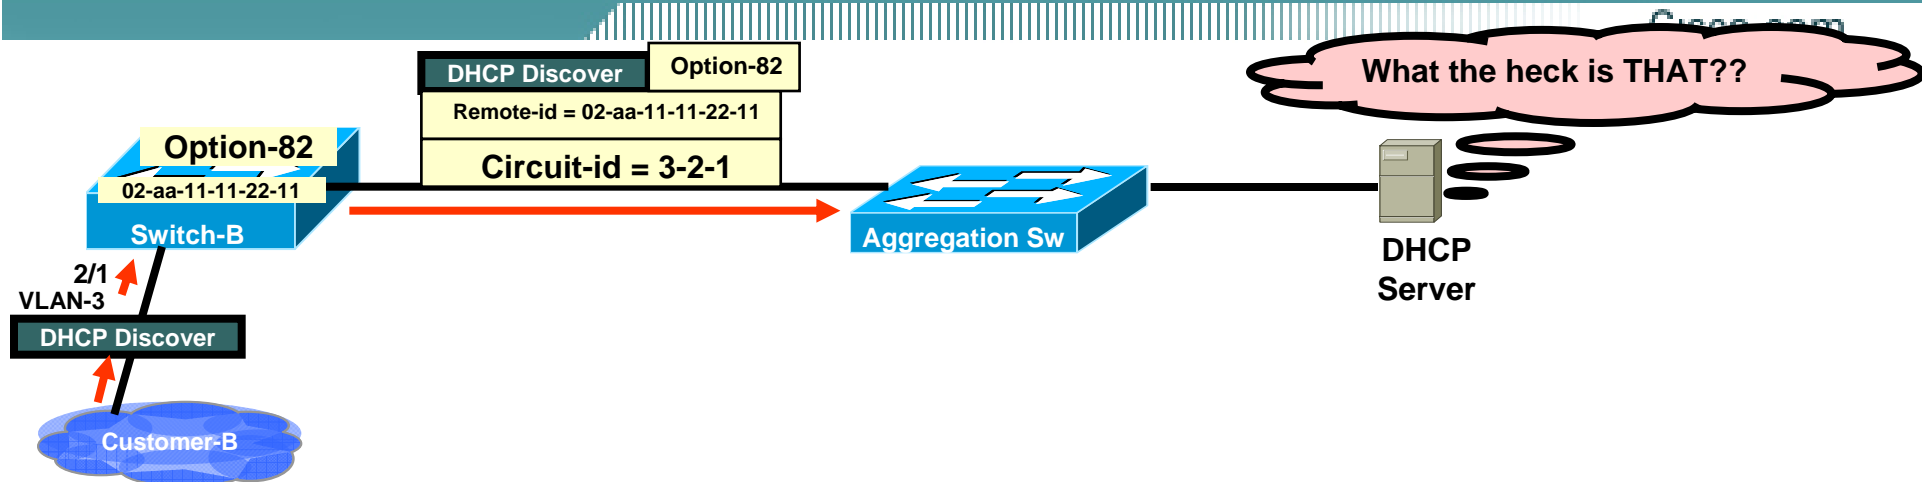
Remote ID Suboption Frame Format



Cat6500# Debug ip dhcp snooping packet


```
Cat6500#
Mar 19 12:33:49.023: dhcp_snooping_draco2_new_dhcp_pak: input idb: Vlan1, vlan 1
Mar 19 12:33:49.023: dhcp_snooping_draco2_new_dhcp_pak:slot = 3, port = 6
Mar 19 12:33:49.023: dhcp_snooping_draco2_new_dhcp_pak:ingress DHCP packet, fixed idb = FastEthernet3/6
Mar 19 12:33:49.023: DHCP_SNOOPING: received new DHCP packet from input interface (FastEthernet3/6)
Mar 19 12:33:49.023: DHCP_SNOOPING: process new DHCP packet, message type: DHCPDISCOVER, input interface: Fa3/6, MAC da: ffff.ffff.ffff, MAC sa: 0011.5c16.4f60, IP da: 255.255.255.255, IP sa: 0.0.0.0, DHCP ciaddr: 0.0.0.0, DHCP yiaddr: 0.0.0.0, DHCP siaddr: 0.0.0.0, DHCP giaddr: 0.0.0.0, DHCP chaddr: 0011.5c16.4f60
Mar 19 12:33:49.023: DHCP_SNOOPING: add relay information option.
Mar 19 12:33:49.023: DHCP_SNOOPING_SW: Encoding opt82 in vlan-mod-port format
Mar 19 12:33:49.023: DHCP_SNOOPING: binary dump of relay info option, length: 20 data:
0x52 0x12 0x01 0x06 0x00 0x04 0x00 0x01 0x03 0x06 0x02 0x08 0x00 0x06 0x00 0x13 0x5F 0x1D 0x7F 0x80
```


DHCP Option-82 Caveats



1. DHCP Servers must be configured to recognize and respond in some way to DHCP Option-82 otherwise packets may be dropped.
2. Switches receiving DHCP messages containing Option-82 will **DROP THEM** if received on an untrusted interface!!
 - The solution for aggregation switches:
Switch(config)# ip dhcp snooping information option
Switch(config)# **ip dhcp snooping information option allow-untrusted**


This is the DEFAULT setting. Remove it if unsupported by the DHCP Server.



Dynamic ARP Inspection (DAI)

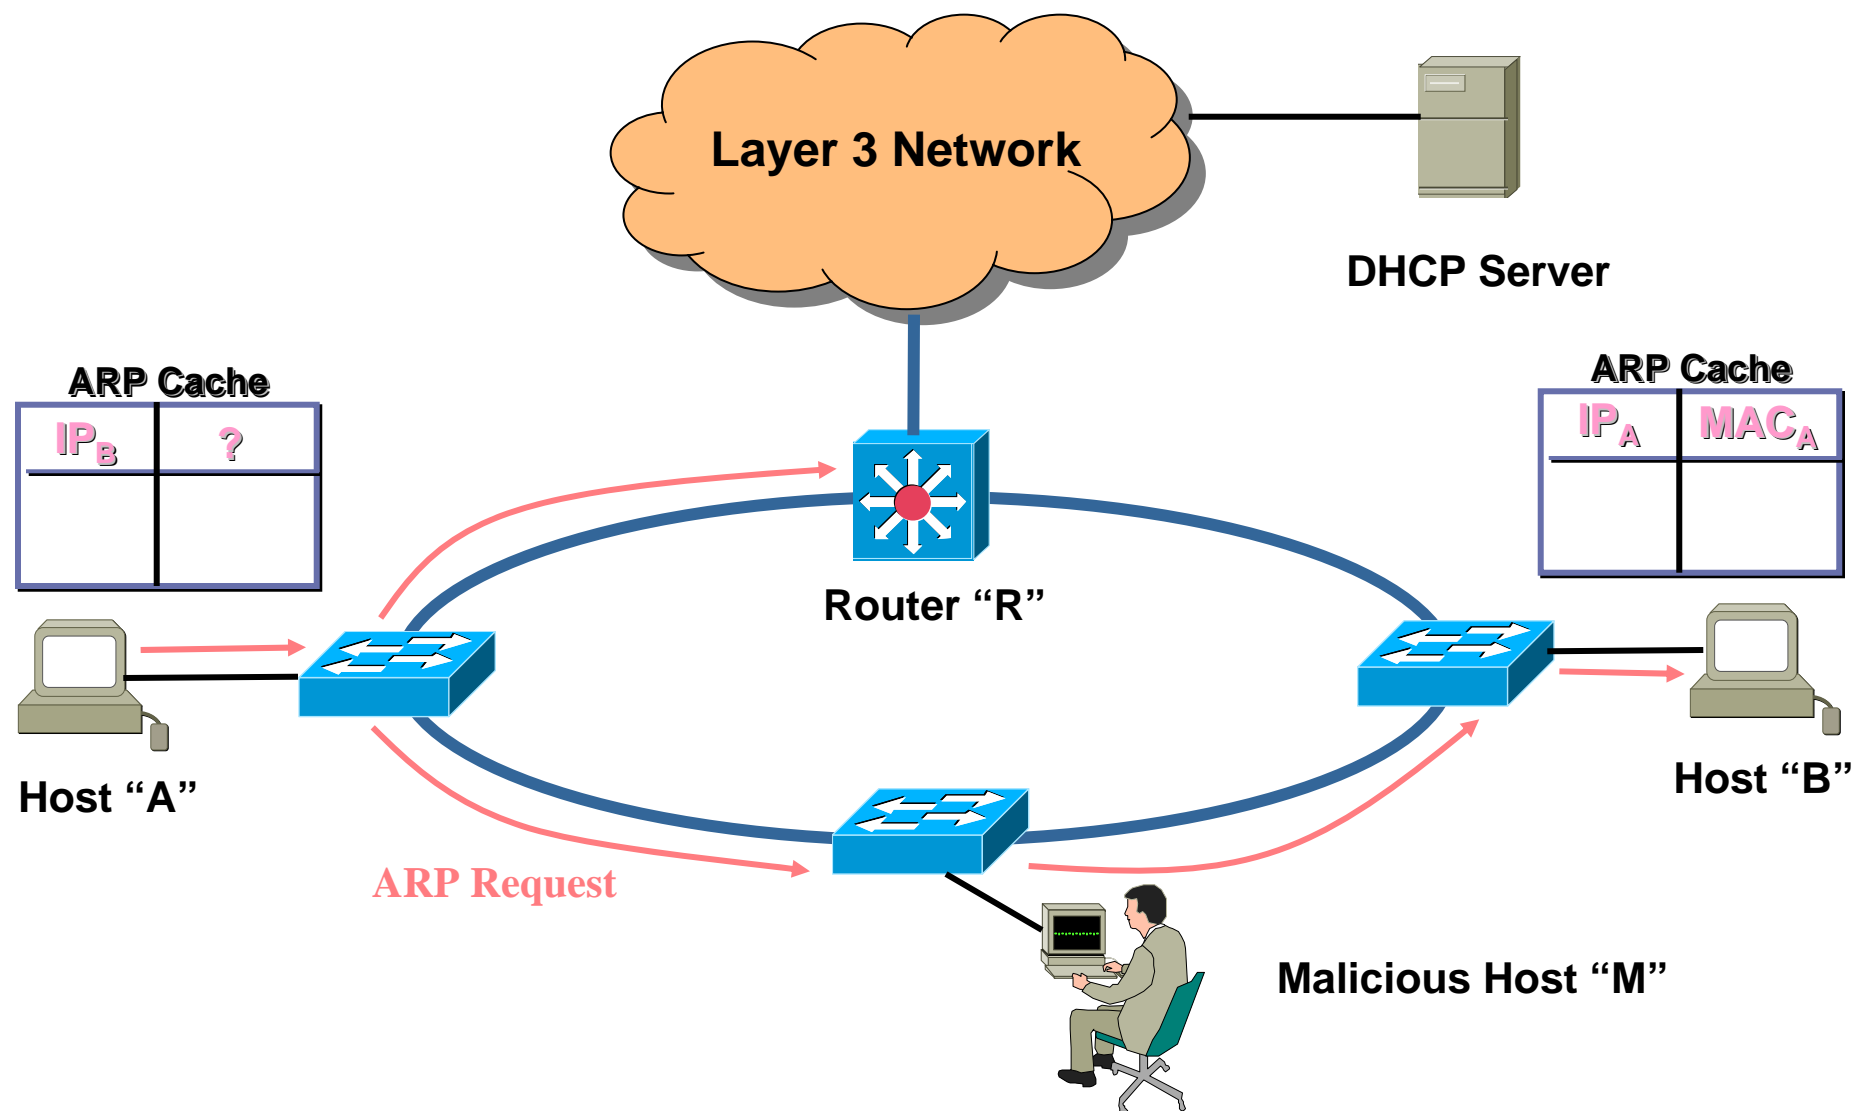
MIM Attack – Attacking another host

Cisco.com



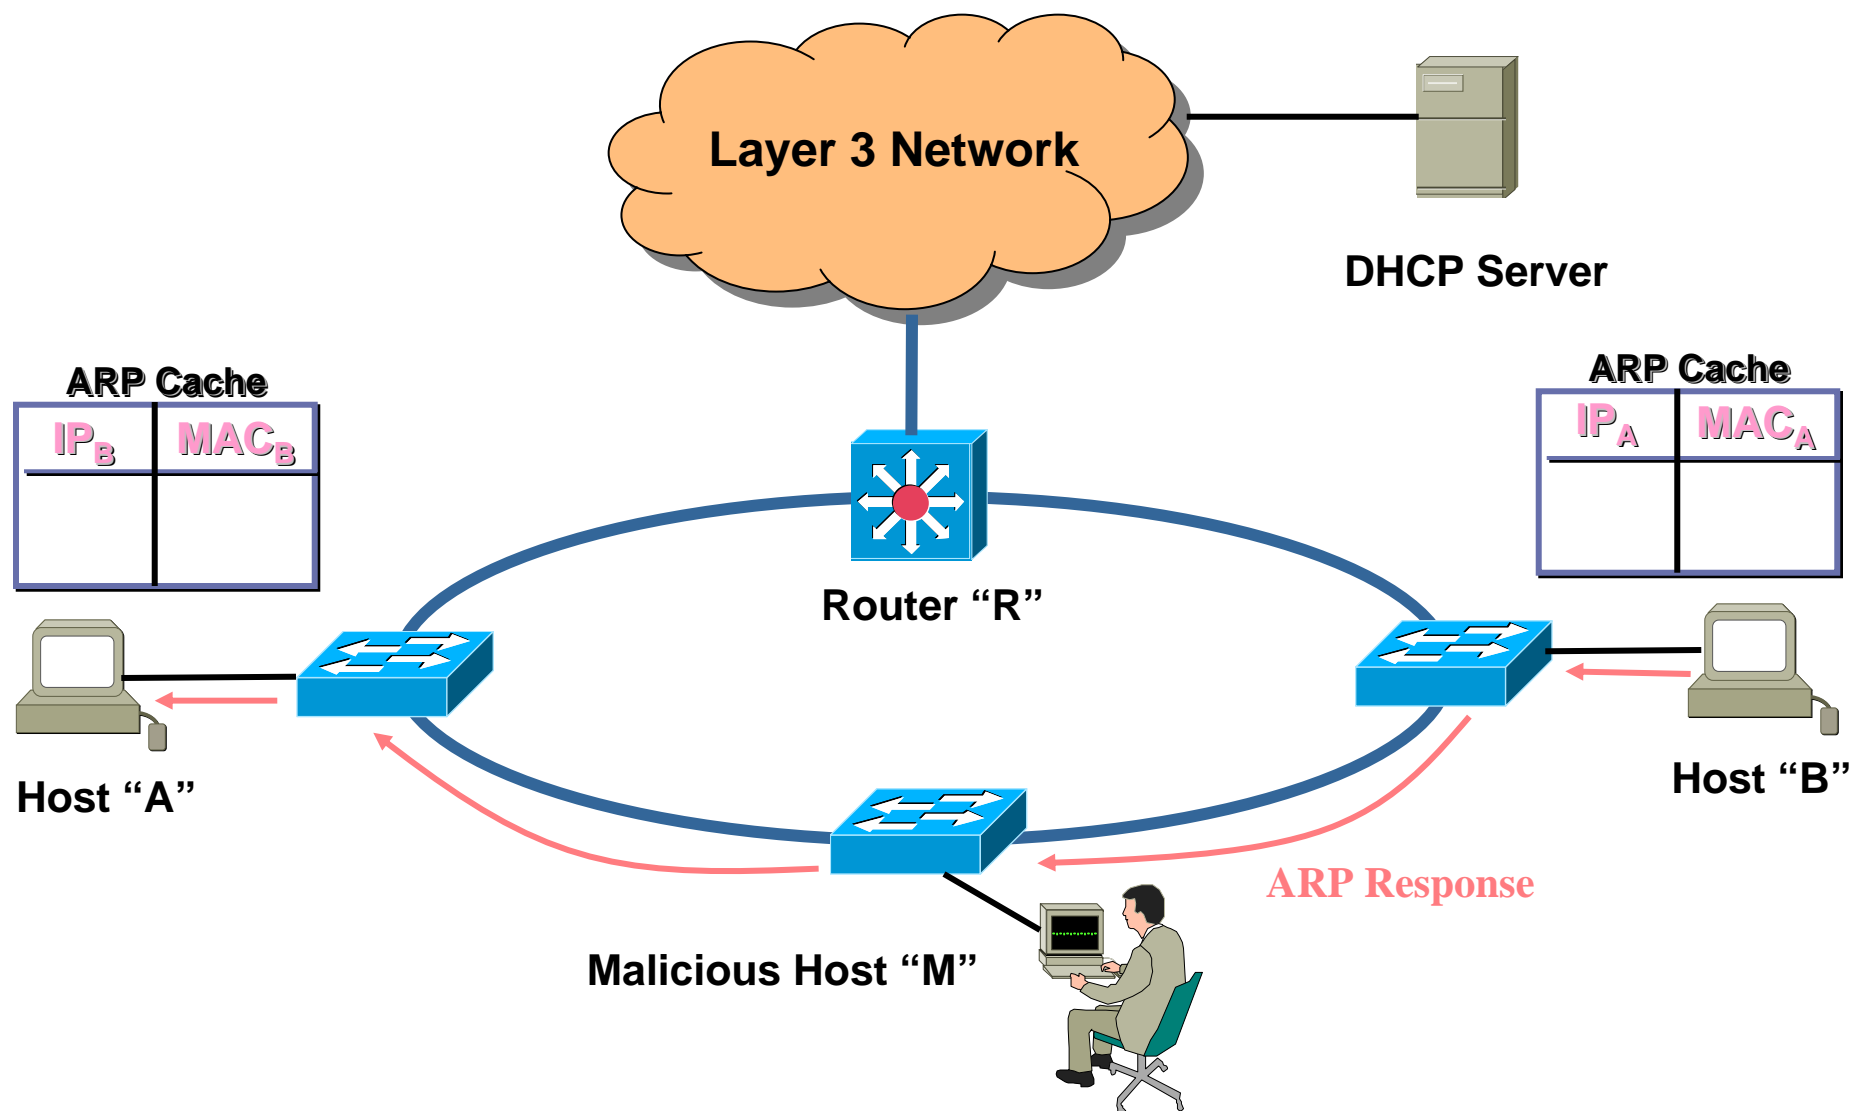
MIM Attack – Attacking another host

Cisco.com




MIM Attack – Attacking another host

Cisco.com



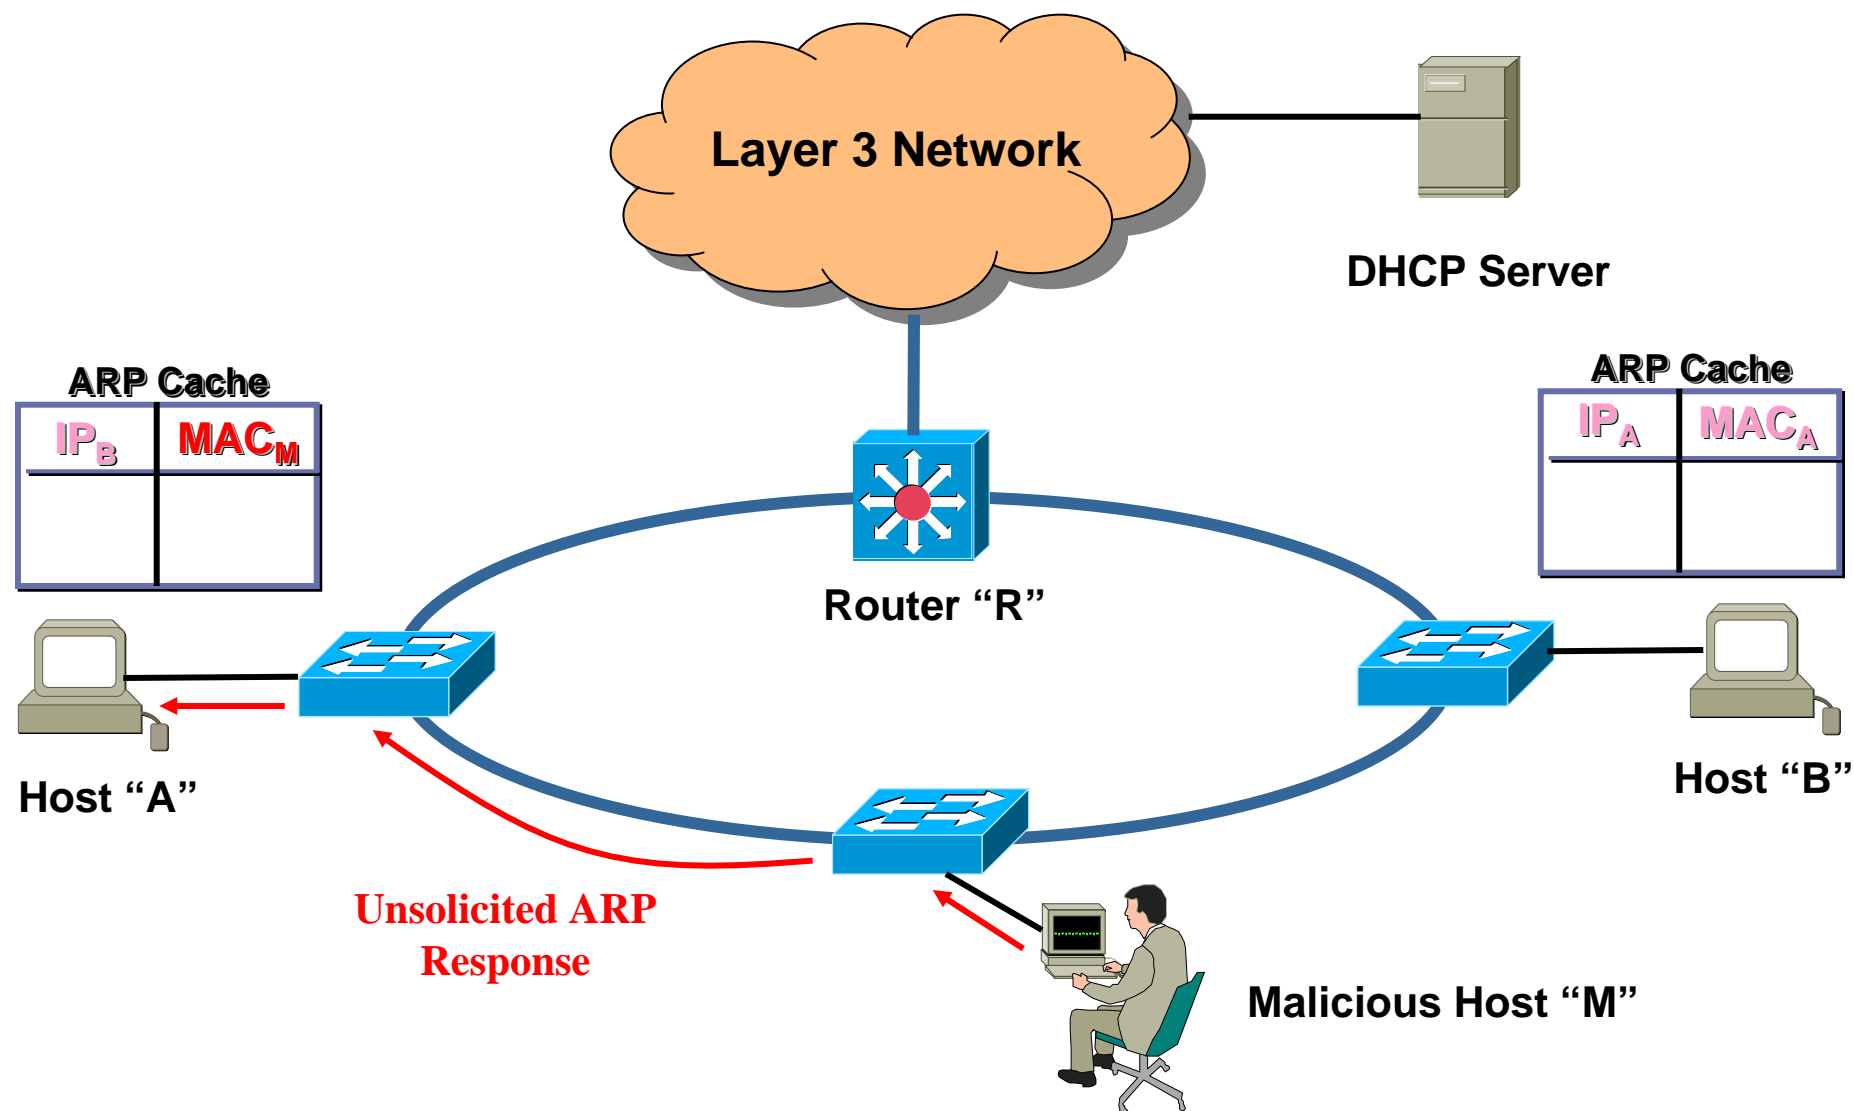
MIM Attack – Attacking another host

Cisco.com



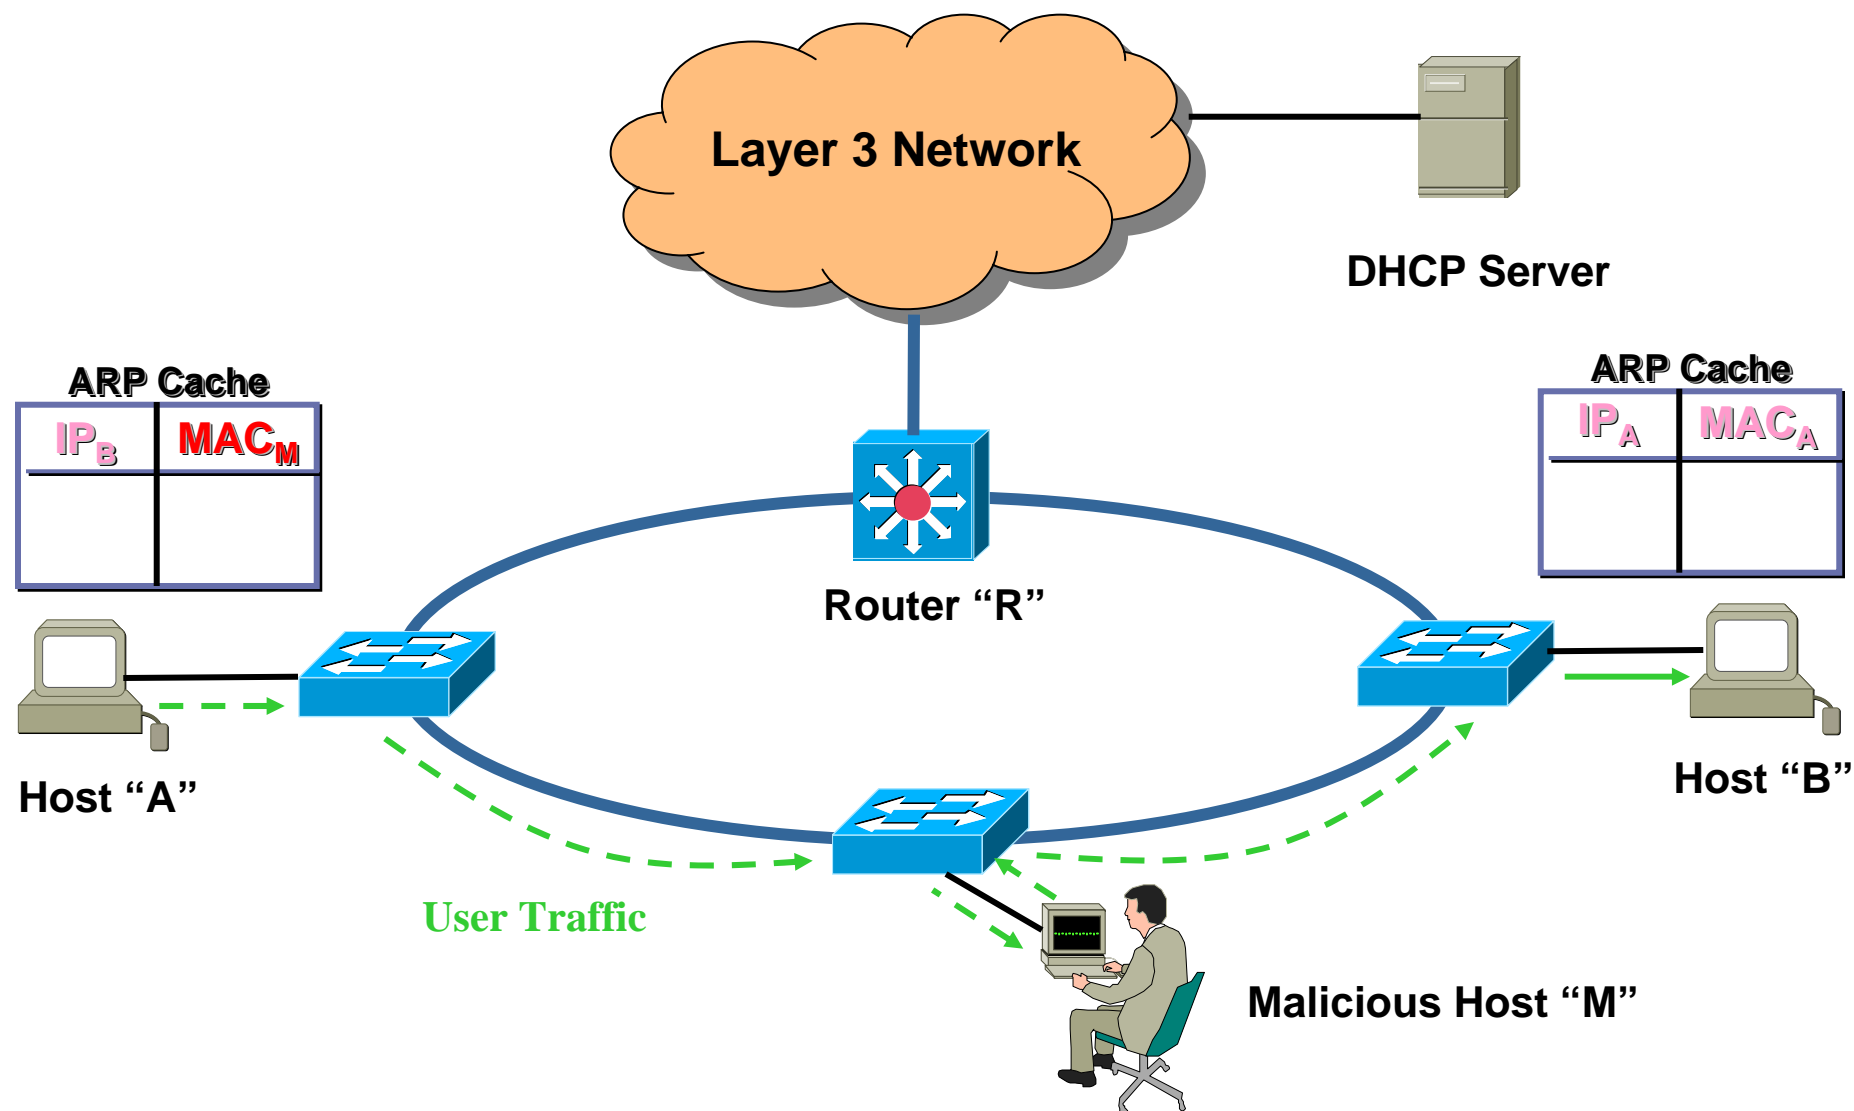
MIM Attack – Attacking another host

Cisco.com



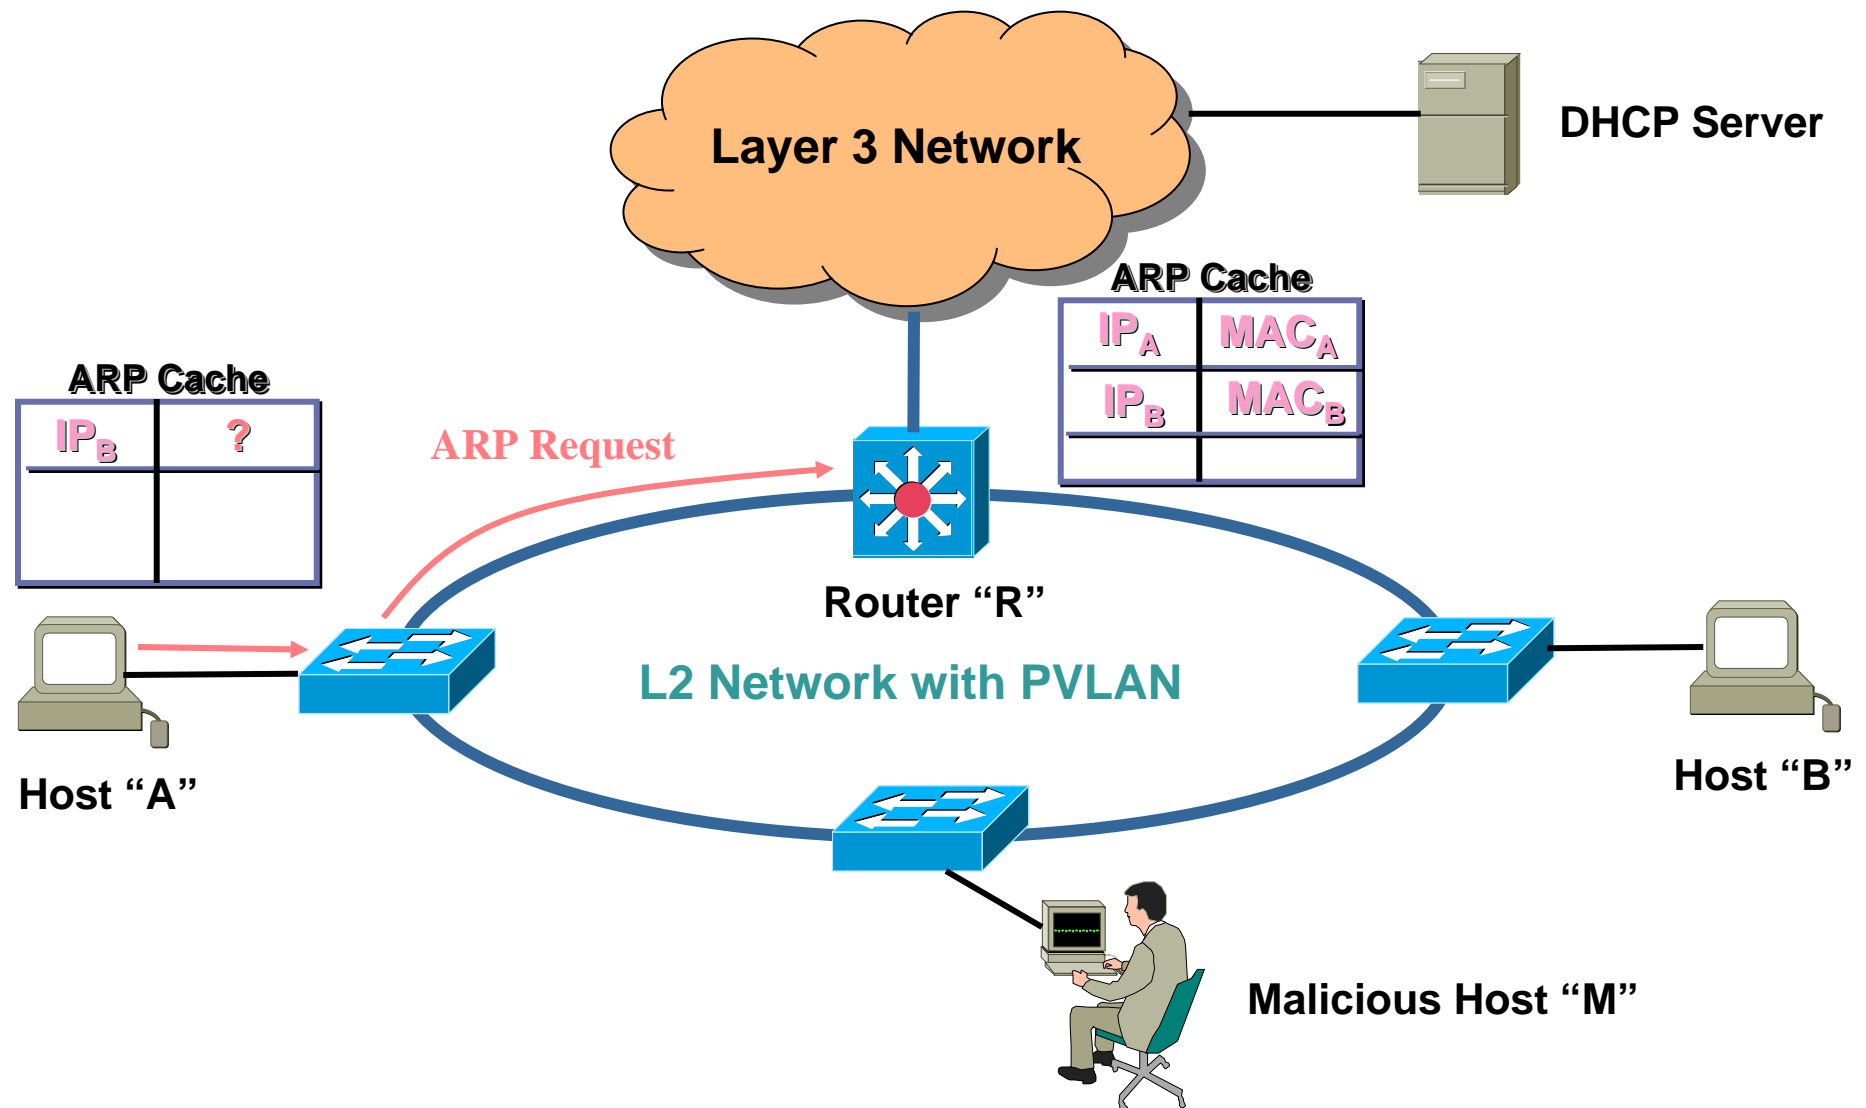
MIM Attack – Attacking another host

Cisco.com



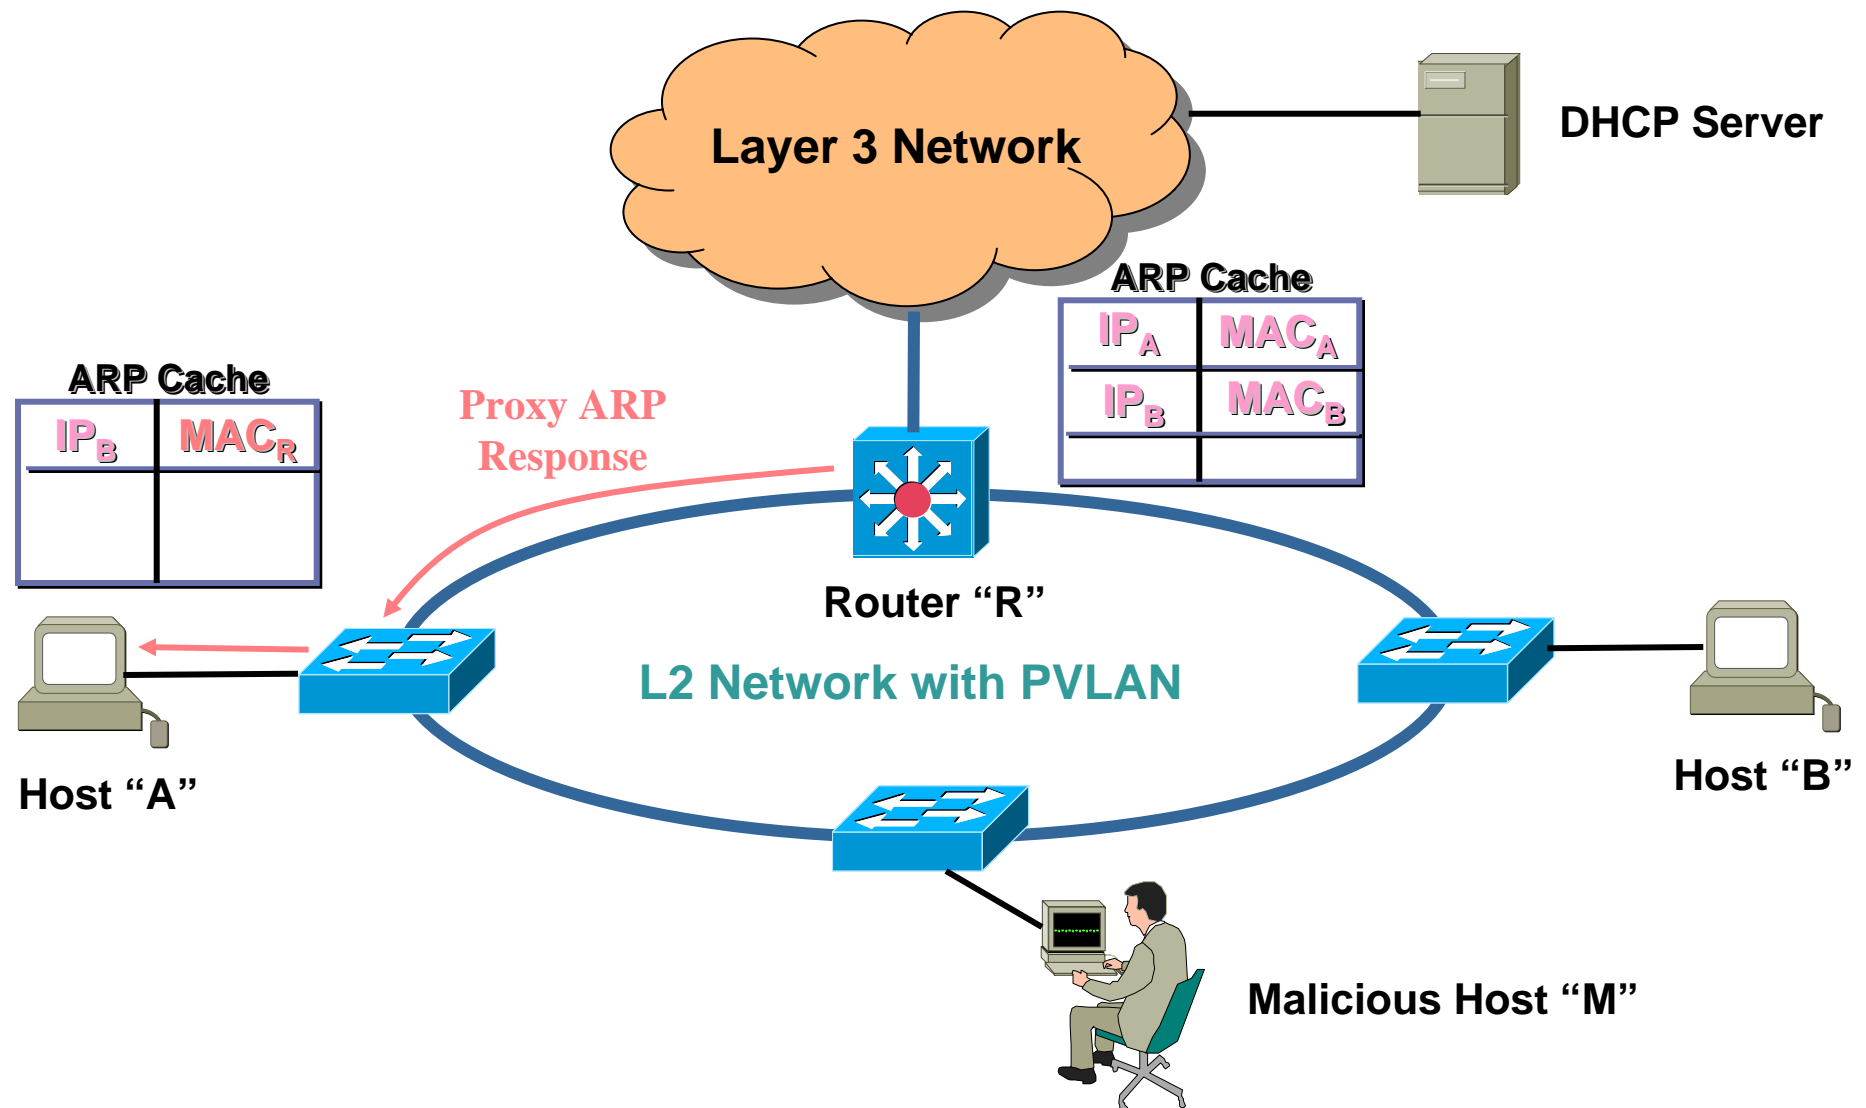
DOS Attack – Attacking the default gateway

Cisco.com



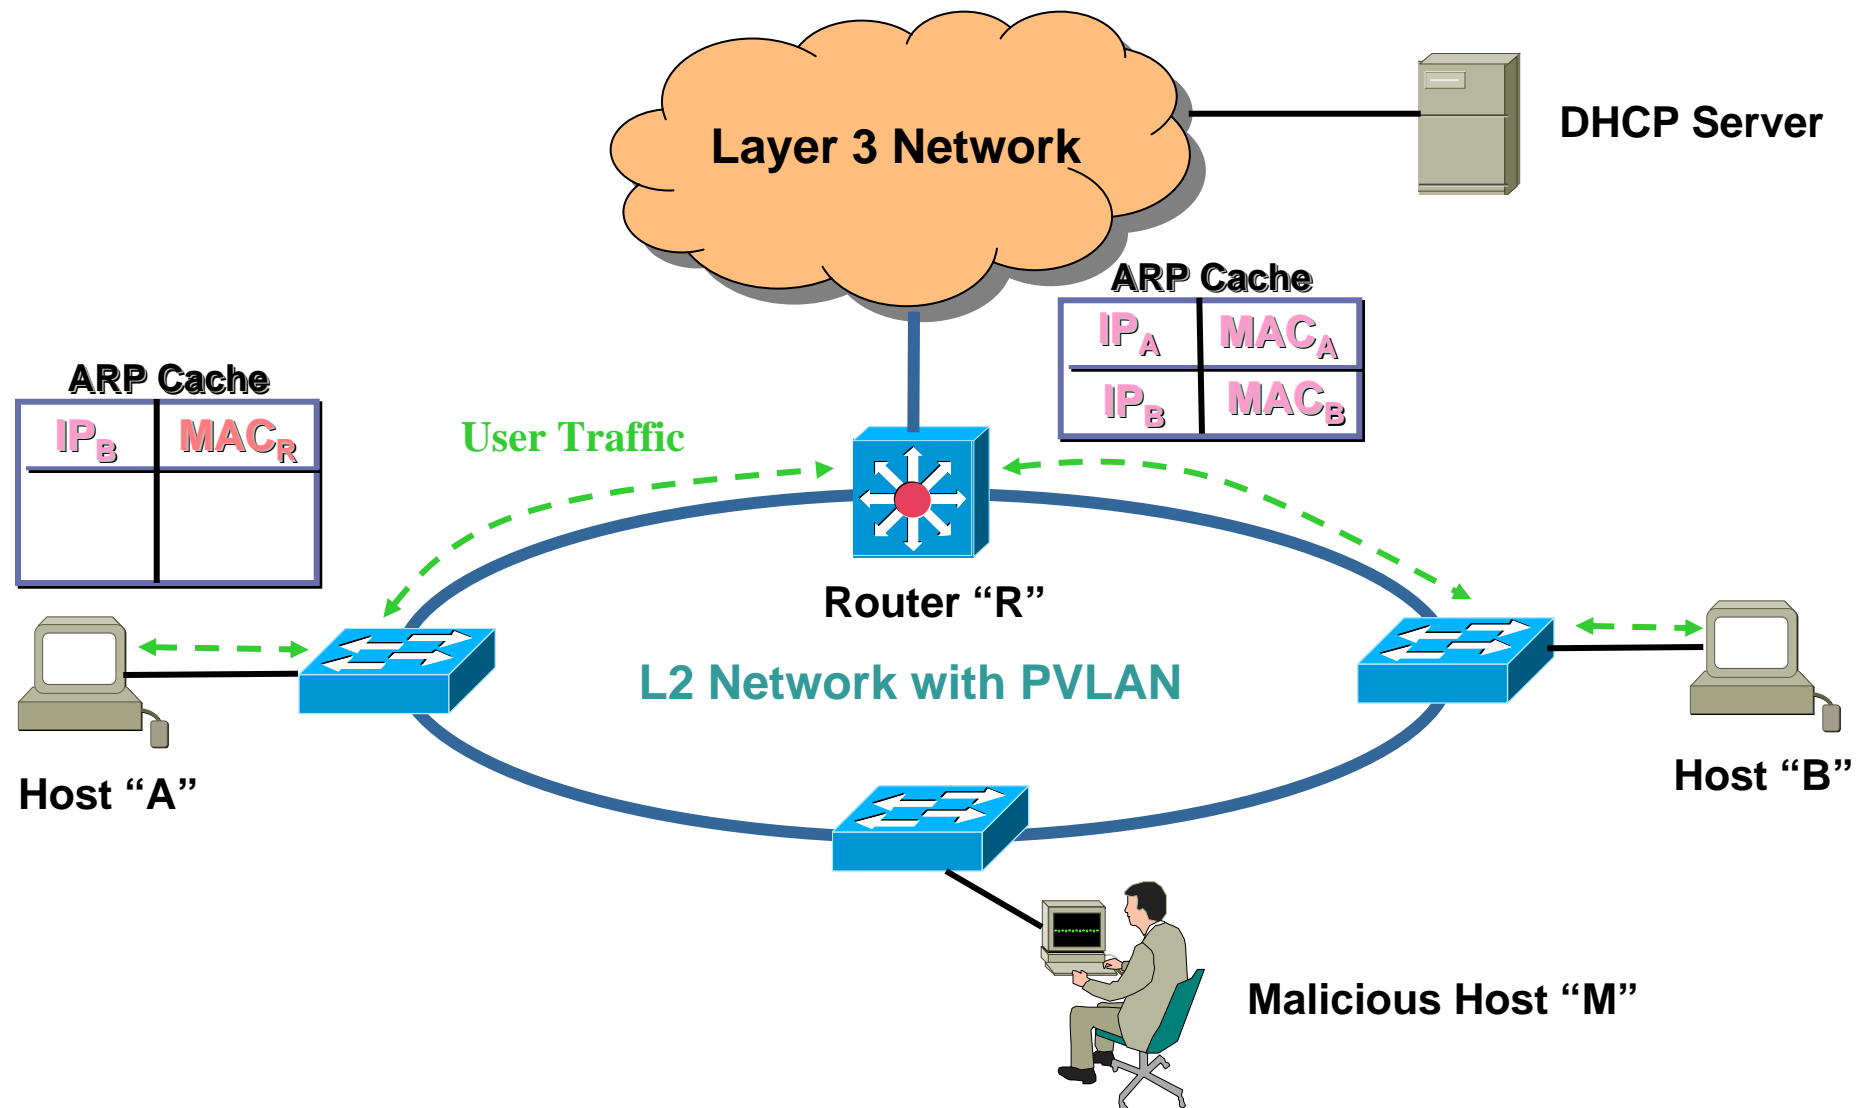
DOS Attack – Attacking the default gateway

Cisco.com



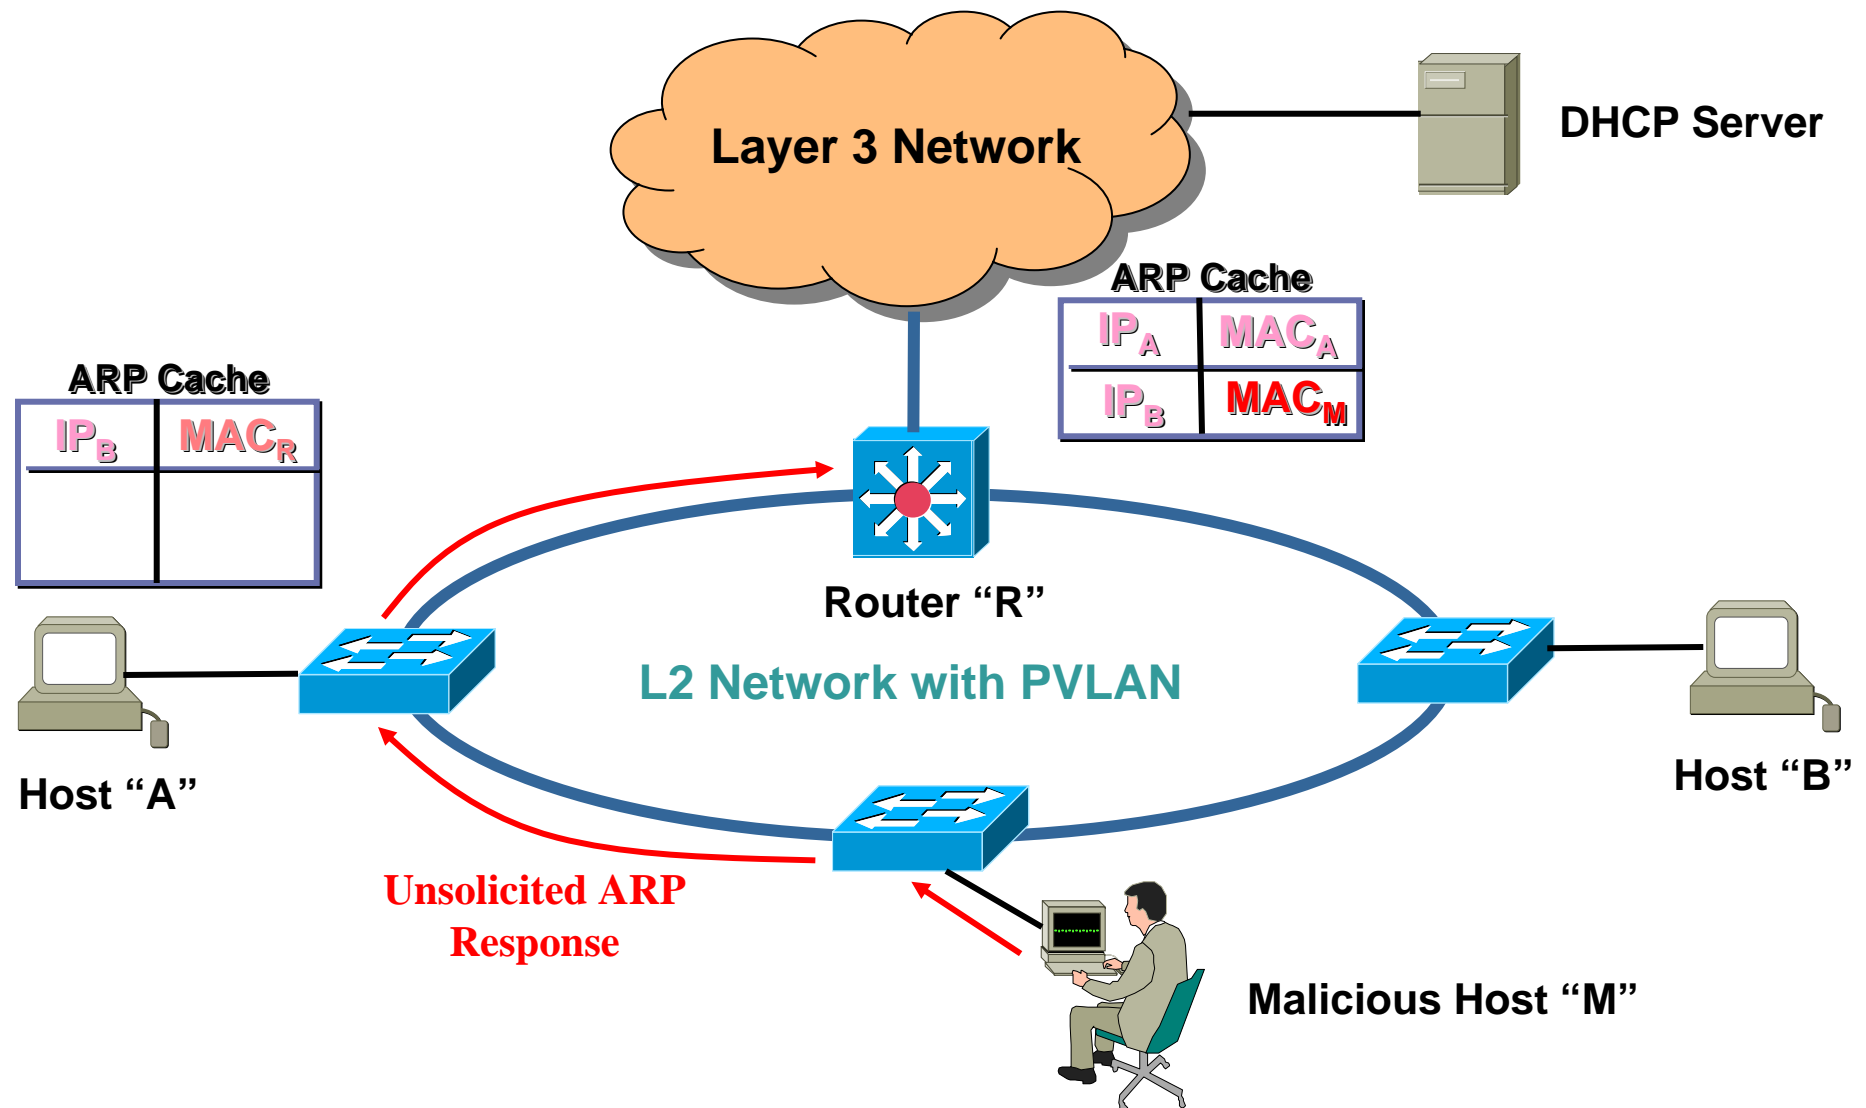
DOS Attack – Attacking the default gateway

Cisco.com



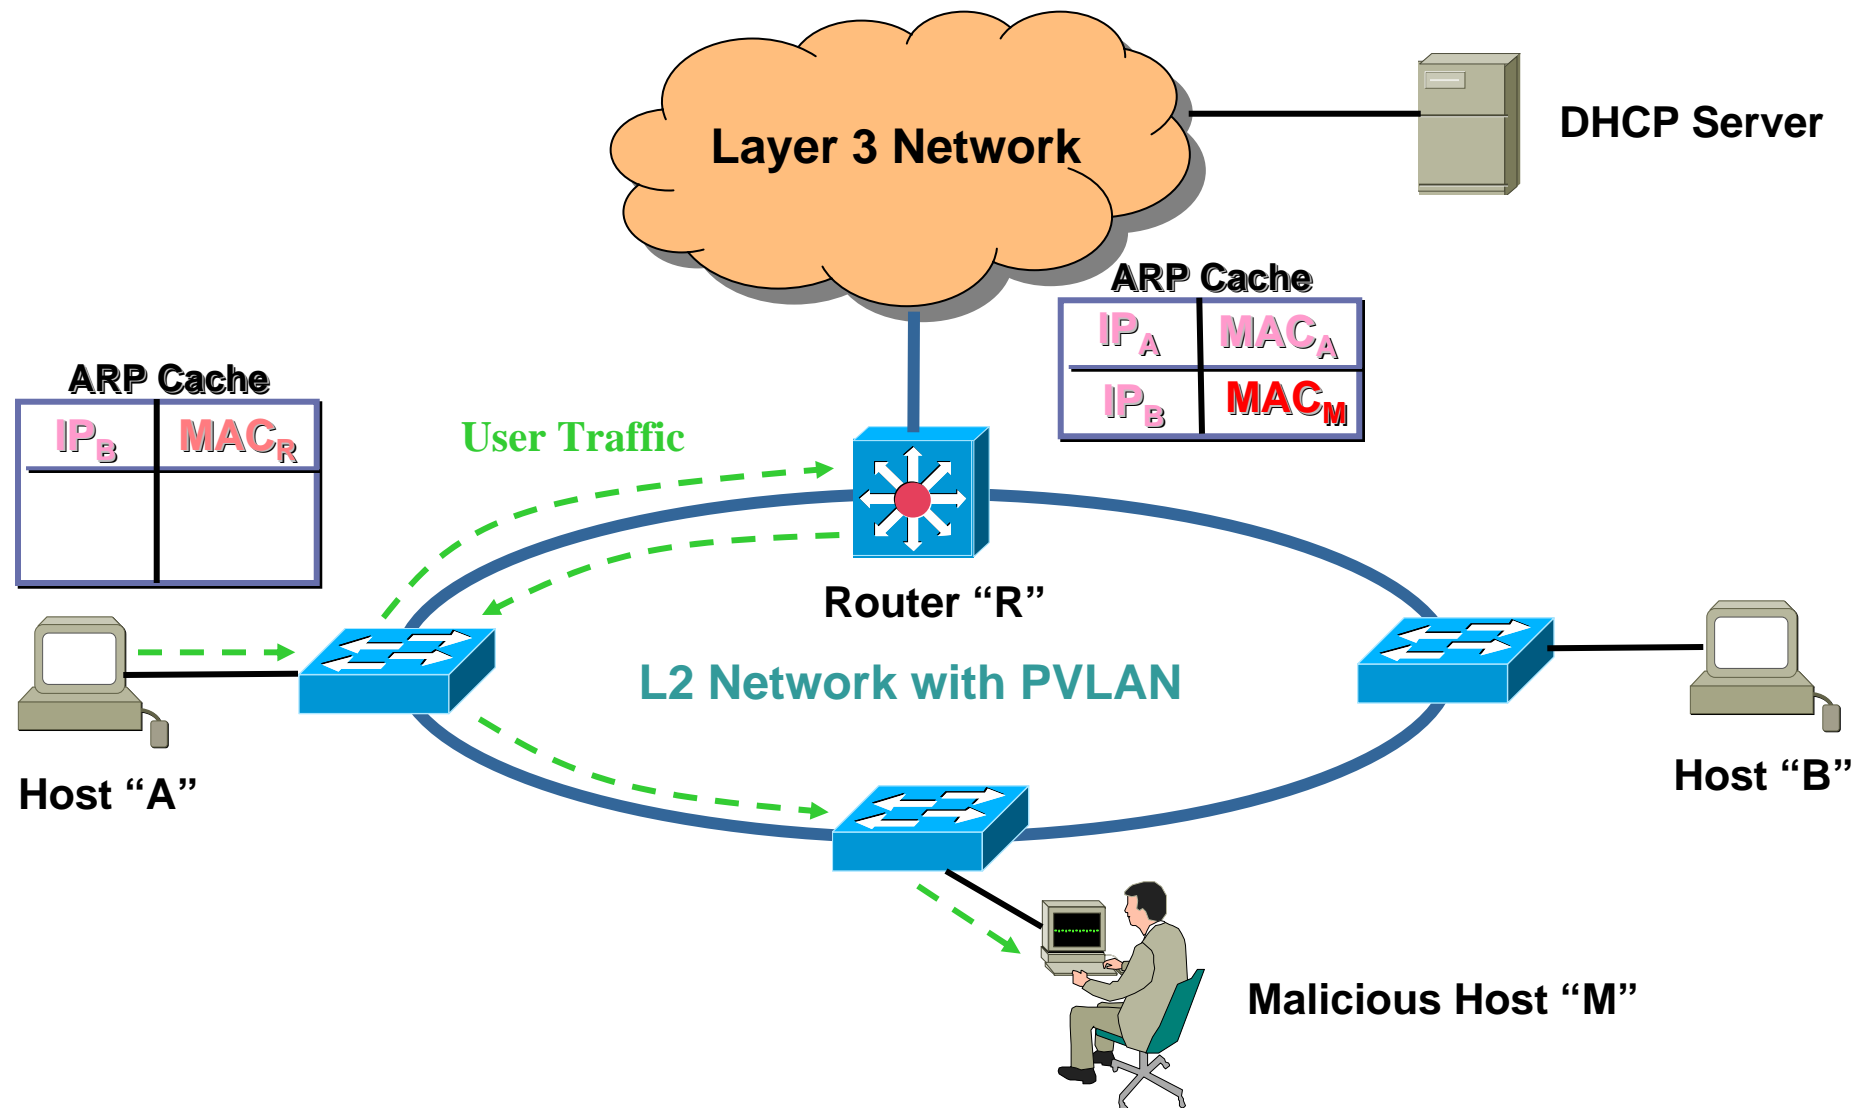
DOS Attack – Attacking the default gateway

Cisco.com



DOS Attack – Attacking the default gateway



Cisco.com



ARP Poisoning: Serious Business

Cisco.com

Recording Voice Calls



- Avaya demonstrated a variation of ARP poisoning at their customer briefing center using Cisco gear
- After intercepting a network connection, packets containing G.711 voice data are collected and the phone conversation is recorded and then replayed
- Demonstrated live to Cisco senior executives in the Cisco network
- Tools are publicly available with GUI and bi-directional spoofs: Ettercap and Dsniff
- Easily taught in 5 minutes
- **Neither the victim nor the default gateway is aware of the attack**


ARP Poisoning Attack Solution: Dynamic ARP Inspection

Cisco.com

Dynamic ARP Inspection – discarding attacker's gratuitous ARP packets in the switch, and logging the attempts for auditing

- Bindings of client IP address, client MAC address, port, VLAN number are built dynamically **by DHCP snooping**
- Switch intercepts all ARP requests and replies on the untrusted access ports
- Each intercepted packet is verified for valid IP-to-MAC binding
- **A solution with no change to the end user or host configurations**

Dynamic ARP Inspection



untrusted

Gratuitous ARP



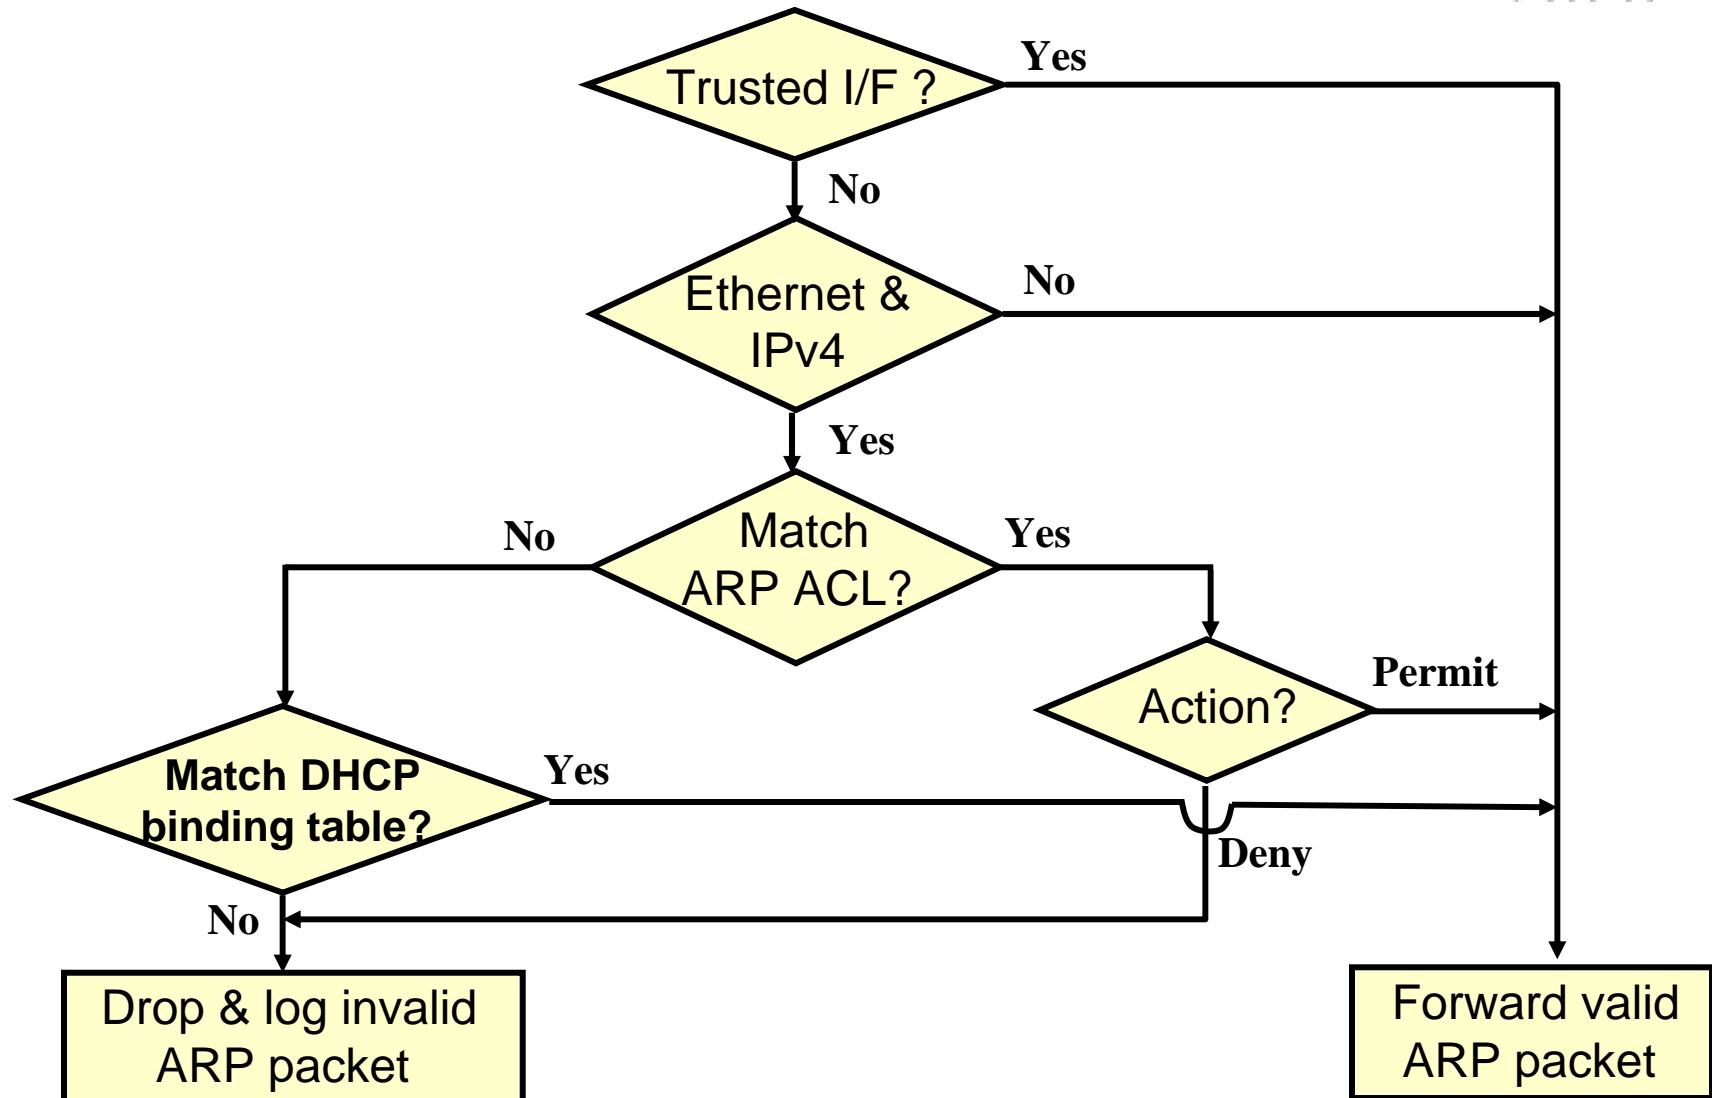
Dynamic ARP Inspection (DAI) Overview

Cisco.com

- **When DHCP Snooping not applicable, static ARP ACLs can be configured instead.**
- **ARP ACLs always take priority over DHCP Snooping Table.**
 - If an ARP ACL is configured to drop a packet, that ARP will be dropped even if there is a valid entry in the DHCP Snooping Table.
- **Relies on same concepts of “Trusted” and “Untrusted” ports as DHCP Snooping.**
 - Ports are untrusted by default
 - **DAI does not verify any ARP Requests/Replies from Trusted interfaces.**

ARP Inspection Procedure

Cisco.com

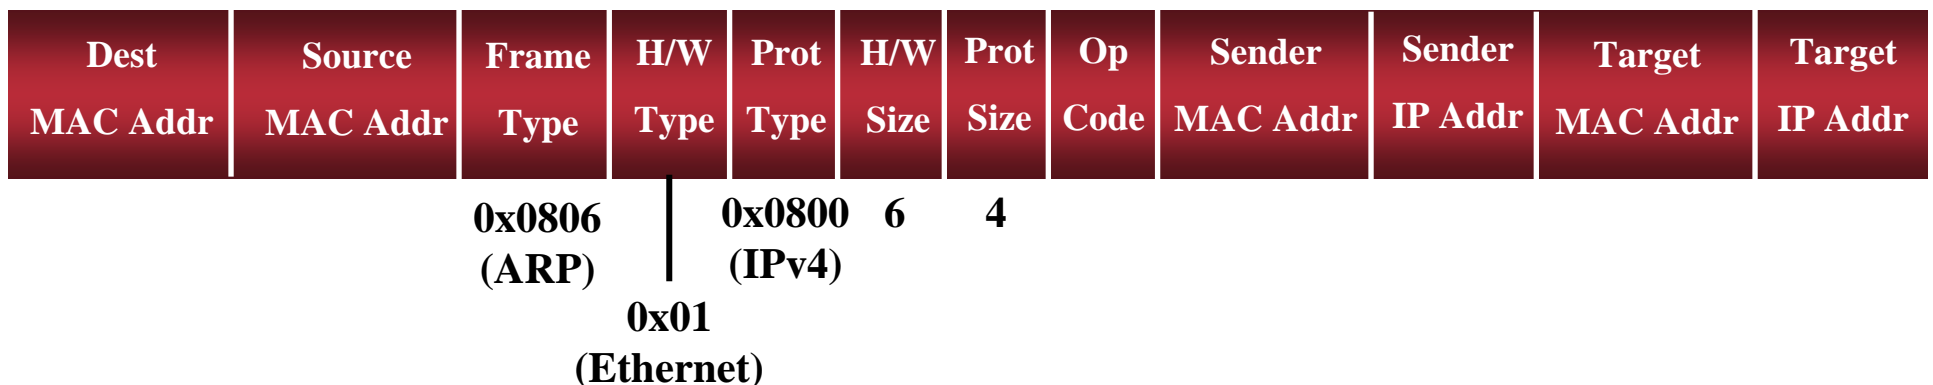


ARP Inspection Overview

Cisco.com

- An ARP request/response packet is considered valid if it meets the following criteria:
 - 1) Mandatory: Sender <MAC, IP, VLAN> triplet is valid
 - 2) Optional: Sender MAC == Source MAC
 - 3) Optional (for ARP response):
Target MAC == Destination MAC

ARP Packet Format



Basic DAI Configuration

Cisco.com

- **Two Design Methodologies:**
 1. **Configure DAI on every switch in the network.**
 - Leave all edge ports as Untrusted
 - Trust all interfaces connected to networking devices (routers, switches, etc).
 2. **Configure DAI on all Edge switches (assuming that hosts are only connected to Edge switches).**
- **Step-1: Configure and verify DHCP Snooping first!**
- **Step-2: Configure DAI:**

```
Cat6500#conf t
```

```
Enter configuration commands, one per line. End with  
CNTL/Z.
```

```
Cat6500(config)#ip arp inspection vlan 1-12
```

```
Cat6500(config)#interface fastethernet3/25
```


```
Cat6500(config-if)#ip arp inspection trust
```

```
Cat6500(config-if)#end
```


DAI in action!! (1)

Cisco.com

DHCP-given
address of 1.1.1.1



3/6
VLAN-1



3/7
VLAN-1

Admin Shut



Fa0/0
1.1.1.1

```
Cat6500#sho ip dhcp snooping bind
```

MacAddress	IpAddress	Lease(sec)	Type	VLAN	Interface
00:11:5C:16:4F:60	1.1.1.1	206234	dhcp-snooping	1	FastEthernet3/6

Total number of bindings: 1

```
Cat6500#sho ip arp inspection
```


Source Mac Validation : Disabled
Destination Mac Validation : Disabled
IP Address Validation : Disabled

Vlan	Configuration	Operation	ACL Match	Static ACL
1	Enabled	Active		

DAI in action!! (2)

Cisco.com

DHCP-given
address of 1.1.1.1



```
Cat6500# sho ip arp inspection interface
```


Interface	Trust State	Rate (pps)	Burst Interval
-----	-----	-----	-----
Gi1/1	Untrusted	15	1
Gi1/2	Untrusted	15	1
Fa3/1	Untrusted	15	1
Fa3/2	Untrusted	15	1
Fa3/3	Untrusted	15	1
Fa3/4	Untrusted	15	1
Fa3/5	Untrusted	15	1
Fa3/6	Untrusted	15	1
Fa3/7	Untrusted	15	1

As soon as the router's FastEthernet interface comes up it will perform a gratuitous ARP...let's see what happens!!


DAI in action!! (3)

Cisco.com

DHCP-given
address of 1.1.1.1



VLAN-1



VLAN-1

Admin Shut



```
Router(config)#int fast 0/0
Router(config-if)#no shutdown
Router(config-if)#^Z
Router#
*Apr 27 02:42:29.805: IP ARP: sent rep src 1.1.1.1 0011.5c10.8340,
                        dst 1.1.1.1 ffff.ffff.ffff FastEthernet0/0
```

```
Router#sh run int fast0/0
!
interface FastEthernet0/0
ip address 1.1.1.1 255.255.255.0
shutdown
duplex auto
speed auto
```

```
Cat6500#
.Mar 20 08:59:28.610: %SW_DAI-4-DHCP_SNOOPING_DENY: 1 Invalid ARPs (Res) on Fa3/7, vlan
1. ([0011.5c10.8340/1.1.1.1/ffff.ffff.ffff/1.1.1.1/08:59:27 PDT Thu Mar 20 2008])
Cat6500#
```

Gratuitous ARP from Router is dropped by DAI on switch.

DAI for non-DHCP hosts

Cisco.com

Notice that in this example, the router has been given a valid, static address of 1.1.1.6 /24. But because it is connected to an untrusted port and does not participate in DHCP, nobody can ARP for it!


```
PC#ping 1.1.1.6

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 1.1.1.6, timeout is 2 seconds:


*Mar 23 04:52:14.175: IP ARP: creating incomplete entry for IP address: 1.1.1.6
interface FastEthernet0/0
*Mar 23 04:52:14.175: IP ARP: sent req src 1.1.1.1 0011.5c16.4f60,
dst 1.1.1.6 0000.0000.0000 FastEthernet0/0.
```

```
Cat6500#
.Mar 20 09:31:37.550: %SW_DAI-4-DHCP_SNOOPING_DENY: 1 Invalid ARPs (Res) on Fa3/7, vlan
1. ([0011.5c10.8340/1.1.1.6/0011.5c16.4f60/1.1.1.1/09:31:37 PDT Thu Mar 20 2008])
```

DHCP-given
address of 1.1.1.1



VLAN-1



VLAN-1

```
Router#sho run int fast 0/0
!
interface FastEthernet0/0
ip address 1.1.1.6 255.255.255.0
duplex auto
speed auto
```



DAI for non-DHCP hosts (2)

Cisco.com

The Solution: ARP Access-List

```
Cat6500#sh run
!
ip arp inspection vlan 1-12
ip arp inspection filter test vlan 1
!
arp access-list test
permit response ip host 1.1.1.6 any mac host 0011.5c10.8340 any
```

Sender of ARP
Response

“any” target
IP address

Senders
MAC address


“any” target
MAC address

```
PC#ping 1.1.1.6
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 1.1.1.6, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/2/4 ms
```

```
Router#show int fast 0/0
FastEthernet0/0 is up, line protocol is up
Hardware is AmdFE, address is 0011.5c10.8340
```

```
Router#sho run int fast 0/0
!
interface FastEthernet0/0
ip address 1.1.1.6 255.255.255.0
duplex auto
speed auto
```

DHCP-given
address of 1.1.1.1



VLAN-1



VLAN-1



ARP ACL Example

Cisco.com

Configuring ARP ACL

```
(Config)#arp access-list arp_acl_1
(config-arp-nacl)# permit ip host 10.1.1.1 mac host 0000.0001.0002
(config-arp-nacl)# deny ip 10.1.1.0 0.0.0.255 mac any
(config-arp-nacl)# permit ip any mac any
```

“IP” will apply to both ARP requests and responses. Alternatively you can also specify “Request” or “Response”.

Applying ARP ACL to a VLAN

```
(config)# ip arp inspection filter arp_acl_1 vlan 5
```

or...

```
(config)# ip arp inspection filter arp_acl_1 vlan 5 static
```

Without the “static” keyword DAI will continue to look for a matching entry in the DHCP Snooping Database if nothing matches the ACL.

With the “static” keyword DAI will use the implicit “deny all” if no match is found in the ACL...even if a corresponding match IS in the DHCP Snooping DB.

Rate-Limiting of ARP traffic

Cisco.com

- ARP packets are rate-limited to prevent a denial-of-service attack on Untrusted interfaces.
- **Default is 15 pps**
- Trusted interfaces are not rate-limited
- (config-if)# **ip arp inspection limit <x>** to raise or lower this limit.
- Exceeding the limit causes the interface to be placed into Errdisable state.

```
Cat6500#  
.Mar 20 16:36:50.180: %SW_DAI-4-PACKET_RATE_EXCEEDED: 16 packets received in 32  
milliseconds on Fa3/7.  
.Mar 20 16:36:50.592: %SW_DAI-4-DHCP_SNOOPING_DENY: 16 Invalid ARPs (Req) on Fa3  
/7, vlan 1. ([0011.5c10.8340/1.1.1.6/0011.5c10.8340/1.1.1.6/16:36:50 PDT Thu Mar  
20 2008])  
Mar 20 16:36:50.184: %PM-SP-4-ERR_DISABLE: arp-inspection error detected on Fa3/  
7, putting Fa3/7 in err-disable state
```

CISCO SYSTEMS



EMPOWERING THE
INTERNET GENERATION